



Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management
- Protection and Security
- Computing Environments
- Open-Source Operating Systems





What is an Operating System?

- A **program** that acts as an intermediary between a **user** of a computer and the **computer hardware**
- **Operating system goals:**
 - (Execute user programs and) make **solving user problems easier**
 - Make the computer system **convenient to use**
 - Use the computer hardware in an efficient manner





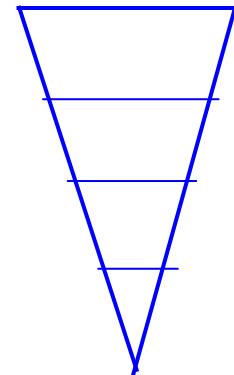
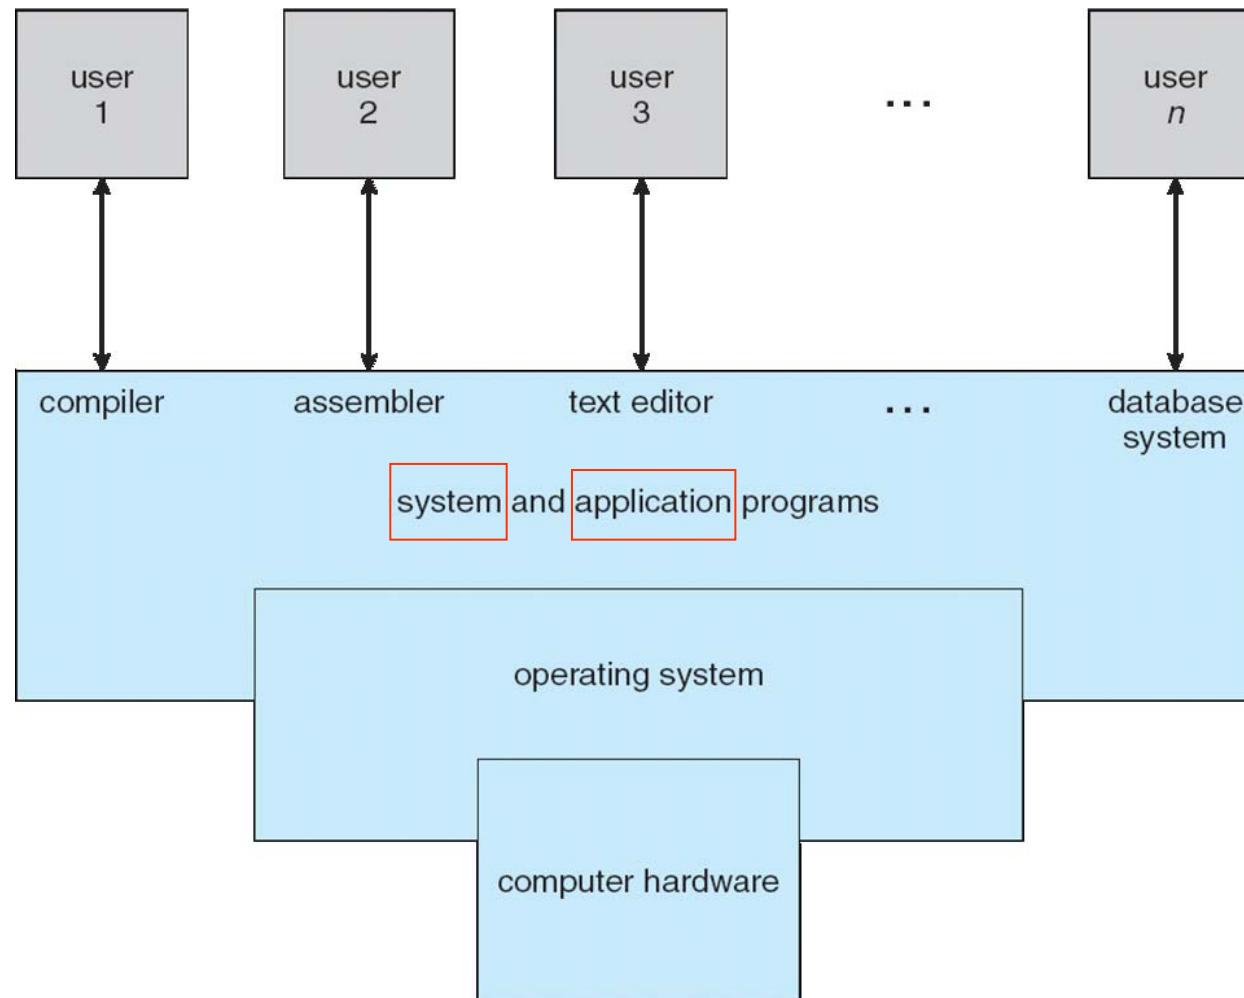
Abstract View of Computer System

- **Computer system** can be divided into four components
 - **Hardware** – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - **Operating system**
 - ▶ Controls and coordinates use of hardware among various applications and users
 - **Application programs** – using the system resources to solve the computing problems of the users
 - ▶ Word processors, web browsers, database systems, video games
 - **Users**
 - ▶ People, machines, other computers





Four Components of a Computer System





Operating System Definition

- OS is a **resource allocator**
 - Manages all resources.
 - Decides between **conflicting requests** for efficient and fair resource use
- OS is a **control program**
 - Control the execution of user programs to **prevent errors** and **improper use** of the computer
- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is good approximation
 - But varies wildly, **Size**: 1MB ~ 1GB
- “The one **program running at all times** on the computer” is the **kernel program**.
 - Everything else is either a **system program** (ships with the operating system) or an **application program**.





Operating System Views

User View:

- Users want convenience, **ease of use**
 - ▶ Don't care about **resource utilization**
- Users of **workstations** have **dedicated resources** but frequently use **shared resources** from **servers**
- **Personal computers**, **text console => GUI**

System View:

- **Mainframe** or **minicomputer** must keep all users happy
- **Handheld computers** (**smart phone**, **tablet**) are **resource poor**, optimized for **usability** and **battery life**
- Some computers have little or **no user interface**, such as **embedded computers** in **devices** and **automobiles**





Computer Startup

● Bootstrap (起動) program

- Executed at **power-up** or **reboot**
- Typically stored in **ROM** or **EPROM**, generally known as **firmware**
- Initializes all aspects of system
 - ▶ check **keyboard**, **main memory**, **hard disk**.
- Loads **operating system kernel** and starts execution

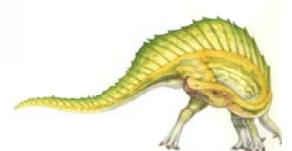
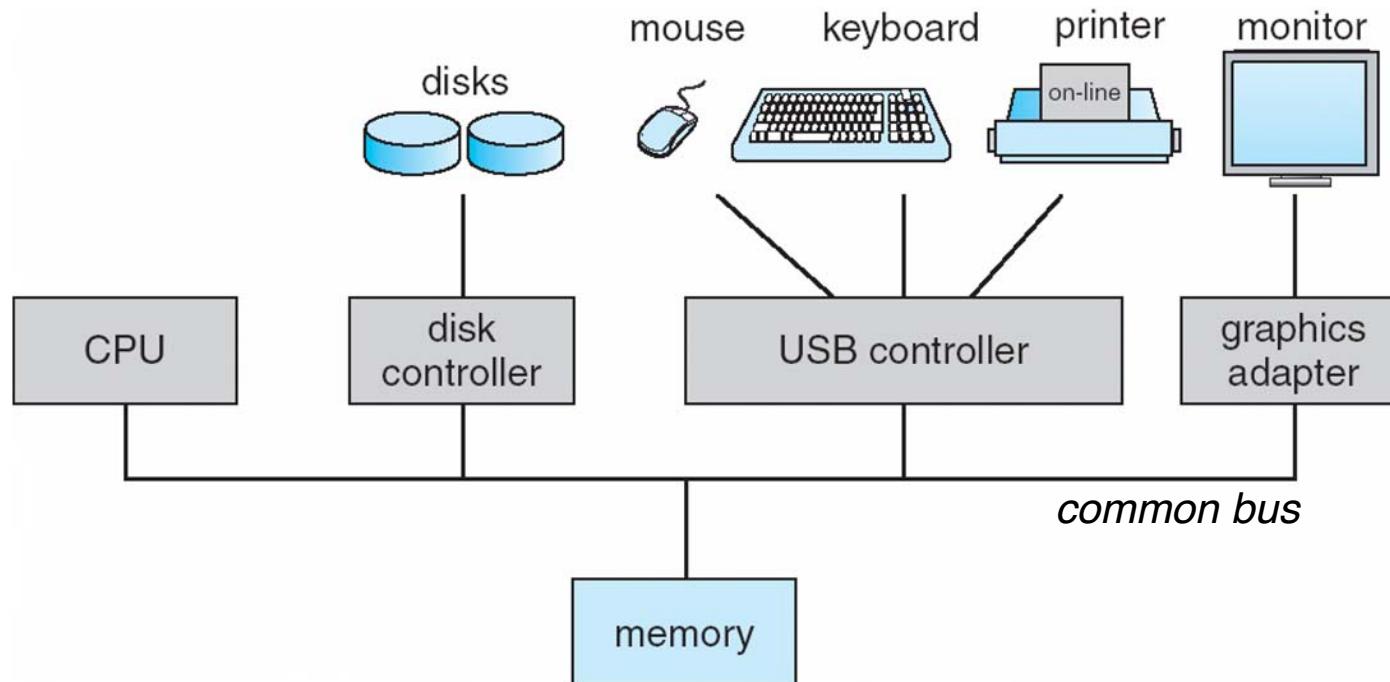




Computer System Organization

■ Computer-system operation (1)

- One or more **CPUs**, **device controllers** connect through **common bus** providing access to **shared memory**
- **Concurrent execution** of **CPUs** and **devices** competing for **memory cycles**





Computer-System Operation

■ Computer-system operation (2)

- I/O devices and the CPU can execute concurrently
 - Each device controller is in charge of a particular device type
 - Each device controller has a local buffer
- I/O devices: keyboard, mouse, printer, hard disk, ...
- Device controller informs CPU that it has finished its operation by causing an *interrupt*
 - CPU moves data from/to main memory to/from local buffers
 - I/O is from the device to local buffer of controller





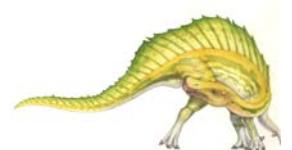
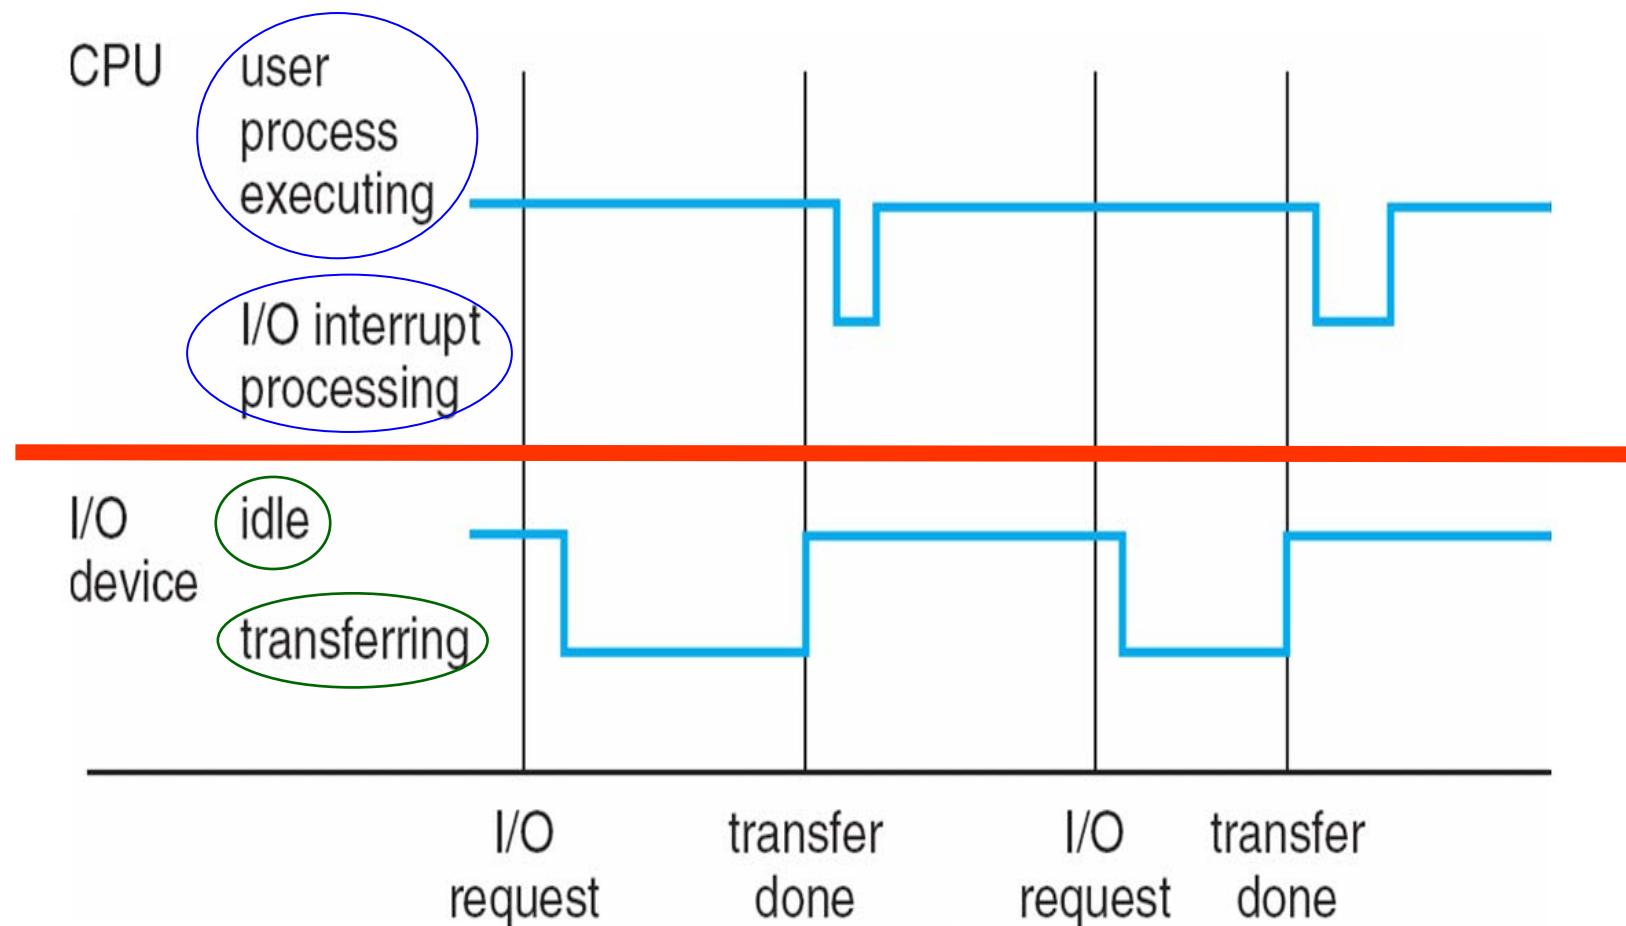
Interrupt Concepts and Handling

- **Interrupt** transfers control to the **interrupt service routine** generally, through the **interrupt vector**, which contains the addresses of all the service routines
- **Interrupt architecture** must save the **address** of the interrupted instruction
- **Incoming interrupts** may be **disabled** while another interrupt is being processed to prevent a **lost interrupt**
- A **trap** is a **software-generated interrupt** caused either by an **error** or a **user request (system call)**
- The **operating system** preserves the state of the **CPU** by storing **registers** and the **program counter (PC)**
- An **operating system** is **interrupt driven**





Interrupt Timeline





I/O Structure

■ Synchronous I/O (e.g. read the keyboard)

- After I/O starts, control returns to user program only upon I/O completion. **No** simultaneous I/O processing
- **System call** – request to the OS to allow a user program to **wait** for I/O completion

■ Asynchronous I/O

- After I/O starts, control returns to user program without waiting for I/O completion. **Can do** simultaneous I/O processing
- **System call** – request to the OS to allow a user program to **not wait** for I/O completion

■ I/O waiting

- **Interrupt**, CPU needs no waiting (modern CPU)
- **Polling**, CPU enter a wait loop and check the status of device





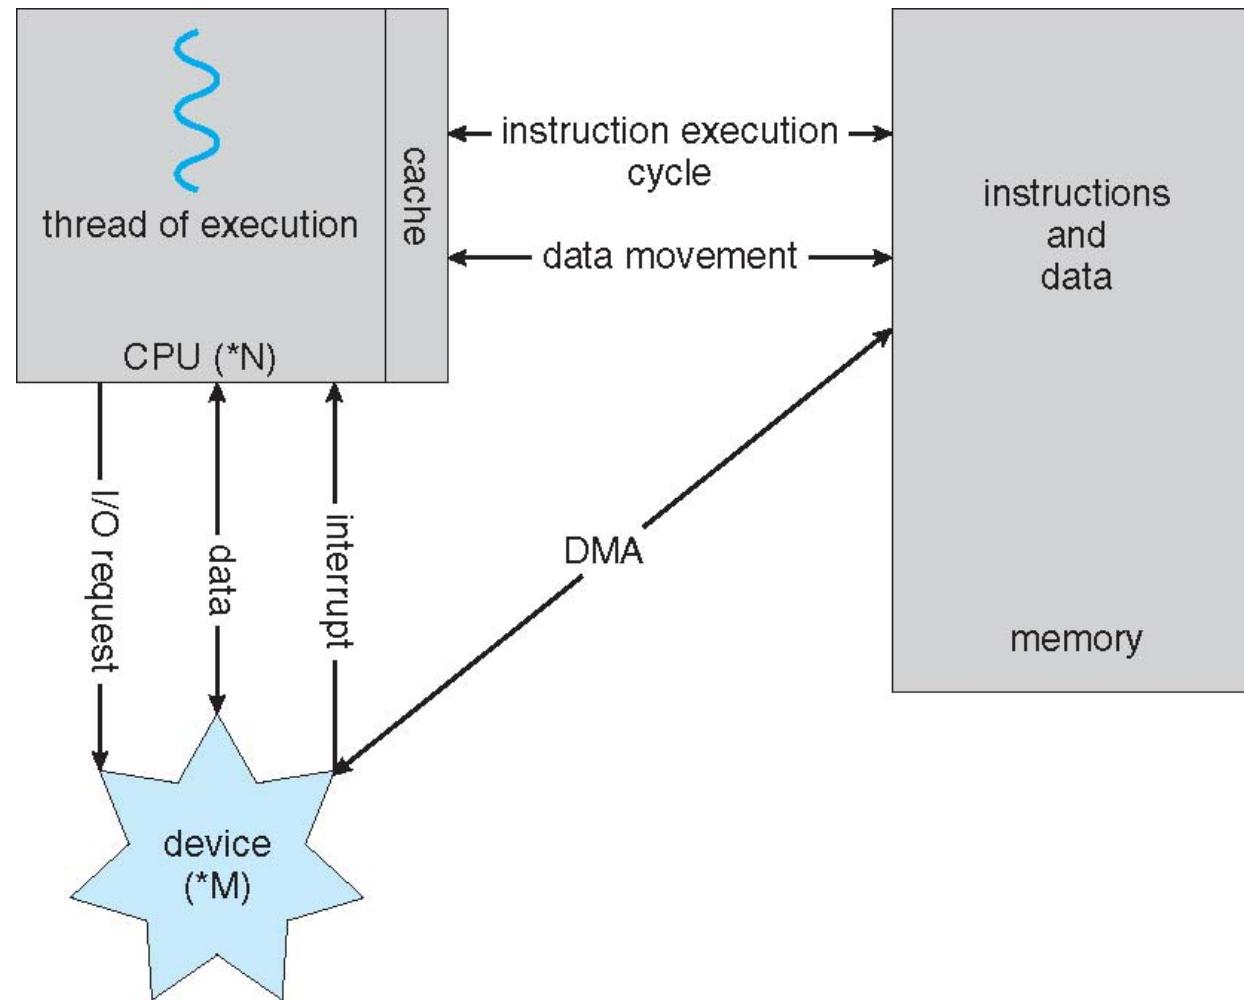
Direct Memory Access Structure

- Only **one interrupt** is generated per **block (block I/O)**, rather than the **one interrupt per byte (character I/O)**
- Used for **high-speed I/O devices** able to transmit information at close to **memory speeds**
- **Device controller** transfers blocks of data from **buffer storage** directly to **main memory** without **CPU intervention**





DMA for Block I/O

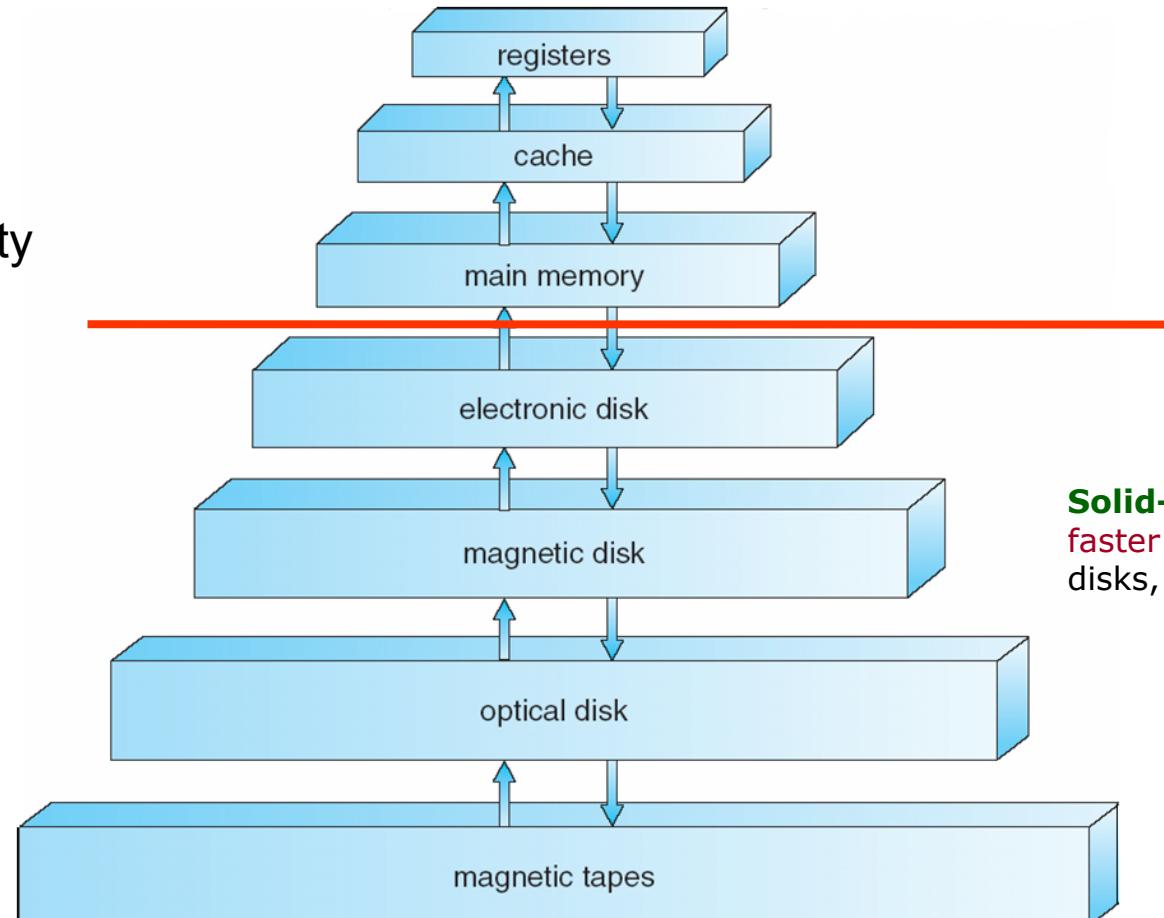




Storage-Device Hierarchy

- Speed
- Cost
- Volatility

tracks
sectors
clusters



Solid-state disks :
faster than magnetic disks, nonvolatile





Caching (暫存)

- Important principle, performed at many levels in a computer (in **hardware**, **operating system**, **software**)
- Copy Information in use from **slower** to **faster** storage temporarily
- **Faster storage** (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache (暫存區) smaller than storage being cached
 - Cache management: important design problem
 - Cache size and replacement policy





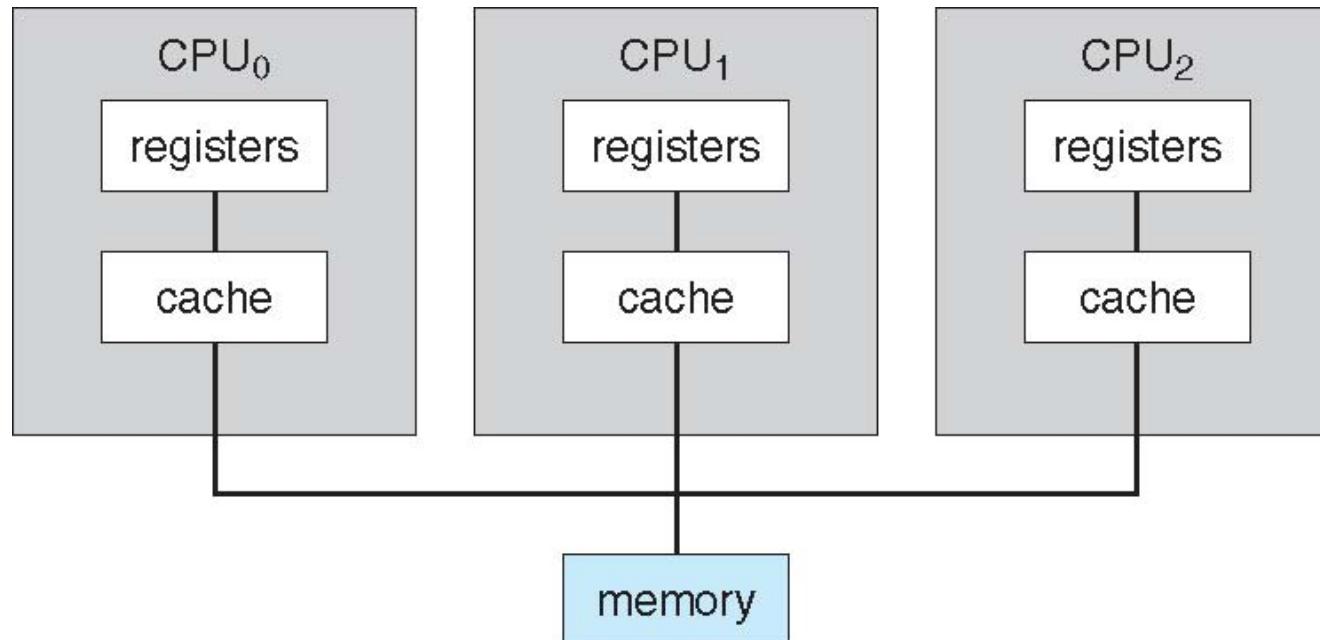
Computer-System Architecture

- Most systems use a **single general-purpose** processor (PDAs through mainframes)
 - Most systems have **special-purpose** processors as well (e.g. GPU)
- **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems, tightly-coupled systems**
 - Advantages include
 1. Increased throughput
 2. Economy of scale
 3. Increased reliability – graceful degradation or fault tolerance
 - Two types in OS design
 1. **Asymmetric Multiprocessing**
 2. **Symmetric Multiprocessing**





Symmetric Multiprocessing Architecture

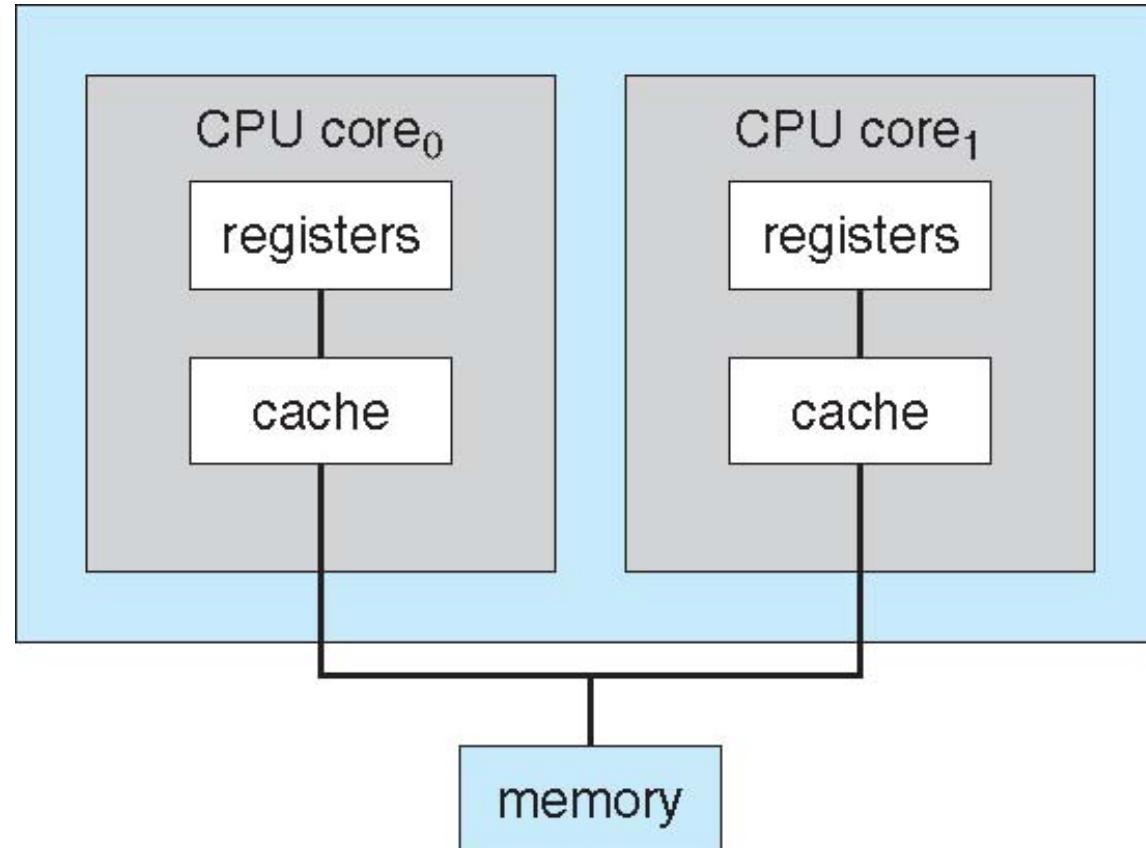


- **UMA** (uniform memory access) and **NUMA** architecture
- Multi-chip and **multicore**





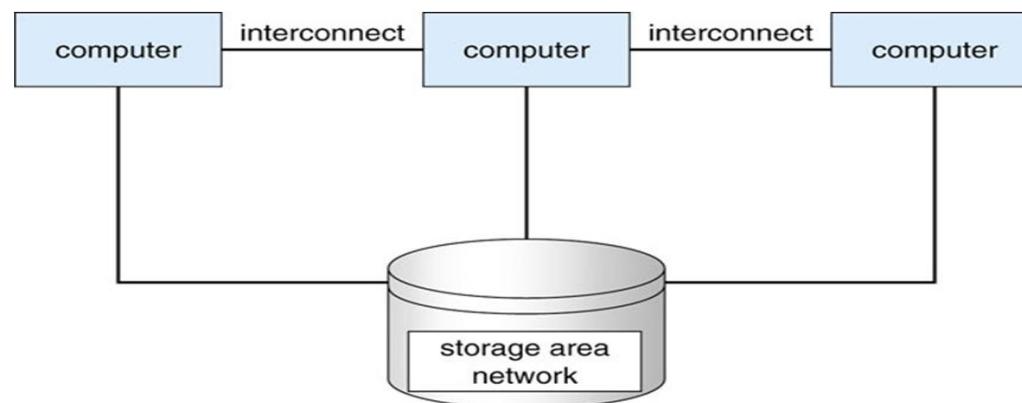
A Dual-Core Design





Clustered Systems (叢集系統)

- Like multiprocessor systems, but multiple systems (OS) working together
 - Usually sharing storage (hard disk) via a storage-area network (SAN, Fig. 1.8)
 - Provides a high-availability service which survives failures
 - ▶ Asymmetric clustering has one machine in hot-standby mode
 - ▶ Symmetric clustering has multiple nodes running applications, monitoring each other
 - Some clusters are for high-performance computing (HPC)
 - ▶ Applications must be written to use parallelization





Operating System Types

■ Batch system

- jobs are **queued** and wait for CPU to execute one at a time
- CPU is **idled** when **waiting I/O** to complete

■ Multiprogramming (多工) OS is needed for efficiency

- Multiprogramming OS organizes jobs (code and data) so CPU always has one to execute
- One job selected and run via **job scheduling**
- When a job has to **wait** (e.g. I/O), **OS switches to another job**

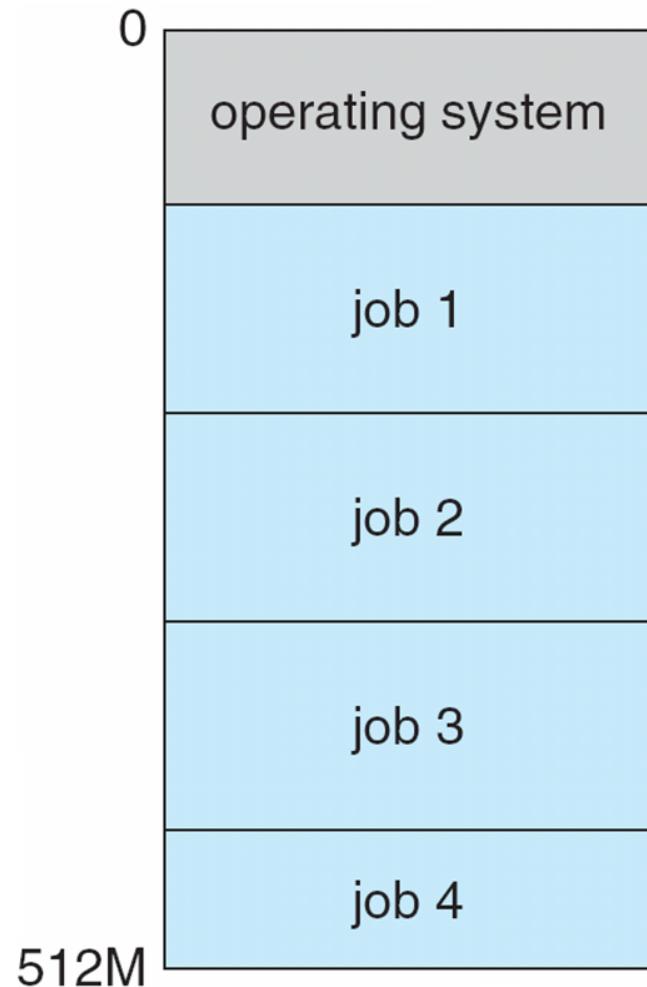
● Timesharing (分時多工, multitasking) OS is a logical extension (=> multiuser)

- CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive computing**
- **Response time** should be < 1 second
- **Each user** has at least one program executing in memory ⇒ **process management**
- If **several jobs ready** to run at the same time ⇒ **CPU scheduling**
- If processes **don't fit in memory**, **swapping** moves them in and out to run





Memory Layout for Multiprogrammed System



Memory management
CPU scheduling

Single-user OS
(e.g. MS Windows)

Multi-user OS
(e.g. Linux)





Operating-System Operations

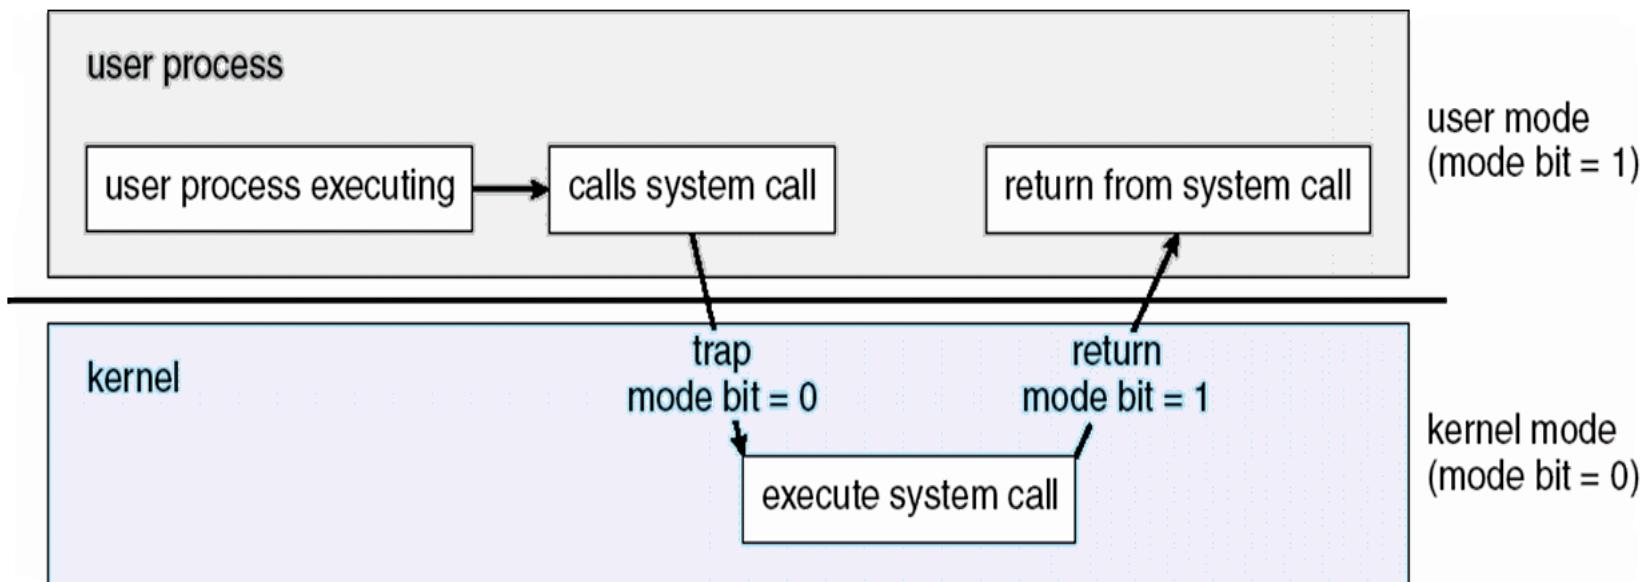
- **Interrupt** driven (interrupts generated by hardware devices)
- **Software error or request** creates **exception** or **trap** (可允許)
 - **division by zero, system call** (request for OS service)
- **Other process problems** include **infinite loop, illegal instruction, illegal memory access** (access other program's memory) (不能允許)
 - require **protection** (CPU and memory protection)
- **Dual-mode** operation allows **OS** to protect itself and other system components (**I/O devices**) (I/O device protection)
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - ▶ to **distinguish** when system is running **user code** or **kernel code**
 - ▶ Some **privileged instructions**, only executable in **kernel mode**
 - ▶ **System call** changes mode to **kernel**, return from call resets it to **user**
 - Intel **8088 CPU** without **mode bit** => **MS-DOS** without protection





Transition from User to Kernel Mode

- **Timer** to prevent **infinite loop** / process hogging resources (**protect CPU**)
 - OS set up a **hardware counter** before scheduling a process to regain **CPU**
 - When **counter come to zero**, an **interrupt** is generated => Enter OS





Process Management - Process

- A **process** is a program in execution.
 - **Program** is a *passive entity*, **process** is an *active entity*.
- **Process** needs **resources** to accomplish its task
 - CPU, memory, I/O, files, Initialization data
- **Process termination** requires **reclaim** of any reusable resources
- **Single-threaded process** has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- **Multi-threaded process** has **one program counter per thread**
- Typically system has **many processes**, some user, some operating system **running concurrently** on one or more CPUs
 - **Concurrency** by multiplexing the CPUs among the processes / threads





Process Management – OS Activities

The **operating system** is responsible for the following activities in connection with **process management**:

- Creating and **deleting** both user and system processes
- **Suspending** and **resuming** processes
- Providing mechanisms for **process synchronization**
- Providing mechanisms for **process communication**
- Providing mechanisms for **deadlock handling**





Memory Management

- All **data** in memory before and after processing
 - All **instructions** in memory in order to execute
 - **Memory management** determines **what** is in memory and **when**
=> Optimizing **CPU utilization** and computer response to users
-
- **Memory management activities**
 - **Keeping track** of which **parts** of memory are currently being used and by **whom**
 - Deciding which **processes** (or parts thereof) and **data** to **move into and out** of memory
 - **Allocating** and **deallocating** memory space as needed





File Management

- OS provides **uniform, logical view** of information storage
 - Storage devices, e.g. disk drive, tape drive
 - ▶ Varying properties include **access speed, capacity, data-transfer rate, access method** (sequential or random)
 - Abstracts physical properties to **logical storage unit** -- **file**
 - Files usually organized into **directories**
 - **Access control** on most systems to determine who can access what
- **OS activities** include
 - Creating and deleting **files** and **directories**
 - Primitives to manipulate **files** and **dirs**
 - Mapping **files** onto secondary storage
 - Backup **files** onto stable (non-volatile) storage media





Mass-Storage Management

- Usually **disks** used to store data that does **not fit in main** memory or data that must be **kept for a long period of time**
- Entire **speed of computer operation** hinges on **disk subsystem** and its algorithms
- **OS activities**
 - **Free-space management**
 - **Storage allocation**
 - **Disk scheduling**
- **Some storage need not be fast**
 - Tertiary (第三級) storage includes **optical storage**, **magnetic tape**
 - Varies between **WORM** (write-once, read-many-times) and **RW** (read-write)
 - Still must be managed





Various Levels of Storage

- Movement between levels of **storage hierarchy** can be explicit or implicit

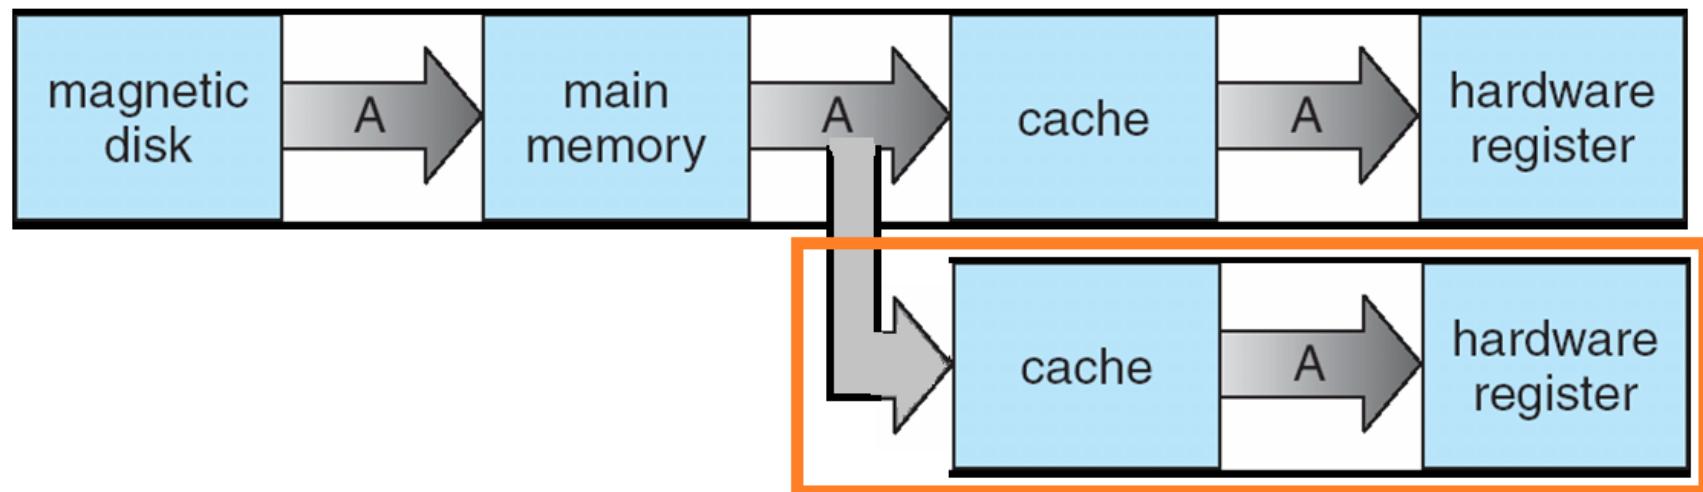
Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000,000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape





Caching: Migration A from Disk to Register

- **Multitasking environments** must be careful to use **most recent value**, no matter where it is stored in the **storage hierarchy**



- **Multiprocessor environment** must provide **cache coherency** in hardware such that all **CPUs** have the most **recent value** in their **cache**
- **Distributed environment** situation even more complex
 - Several copies of a datum can exist





Kernel I/O Subsystem

- One purpose of **OS** is to
 - **hide** peculiarities of **hardware devices** from the **user**
- **I/O subsystem** responsible for
 - **I/O scheduling** (scheduling the I/O requests), **disk r/w requests**
 - **buffering** (storing data temporarily before being transferred), **write file**
 - **caching** (storing data in faster storage for performance), **disk cache**
 - **spooling** (overlapping output of one job with input of other jobs), **printing**
 - General **device-driver interface**
 - **Drivers** for specific **hardware devices**





Protection and Security

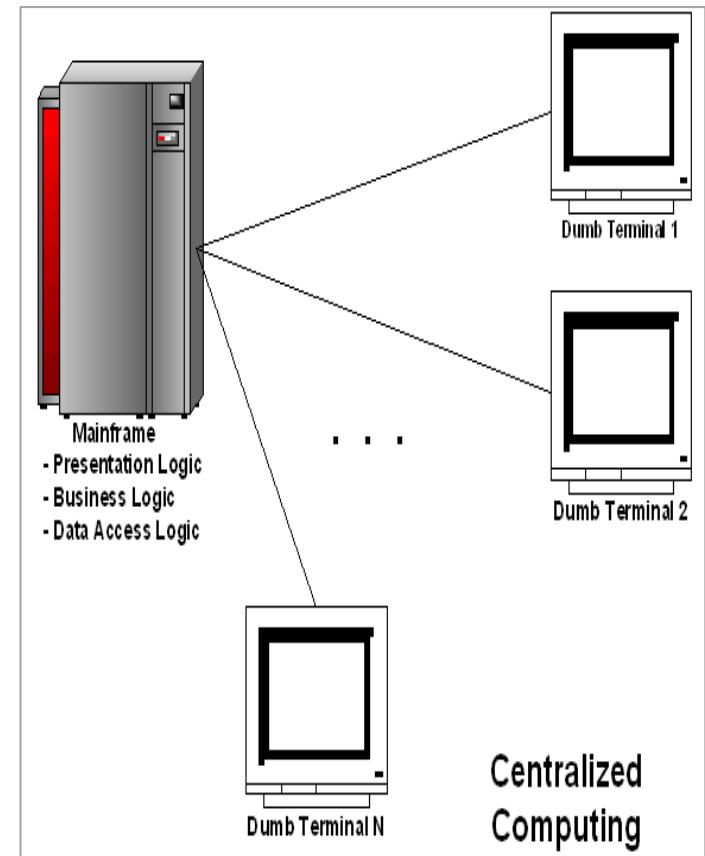
- **Protection** – any mechanism for **controlling access** of **processes** or **users** to **resources** defined by the OS
- **Security** – **defense** of the system against **internal** and **external attacks**
 - Huge range, including **denial-of-service**, **worms**, **viruses**, **identity theft**, theft of service
- **OS** first distinguish among users, to determine **who** can do **what**
 - **User identity** (**UID**, security ID) include **user name** and associated **number**, one per user
 - **UID** then associated with all **files**, **processes** of that user to determine access control
 - **Group identifier** (**GID**) allows **set of users** to be defined. It is also associated with each **process**, **file**
 - **Privilege escalation** (**提升**) allows user to change to **effective ID (set UID)** with more rights





Computing Environments - Traditional

- Stand-alone general purpose machines
 - But blurred as most systems **interconnect** with others (i.e. **the Internet**)
- Portals provide **web access** to internal systems
- Network computers are like **web terminals**
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous –
 - even home systems use **firewalls** to protect home computers from Internet attacks





Computing Environments - Mobile

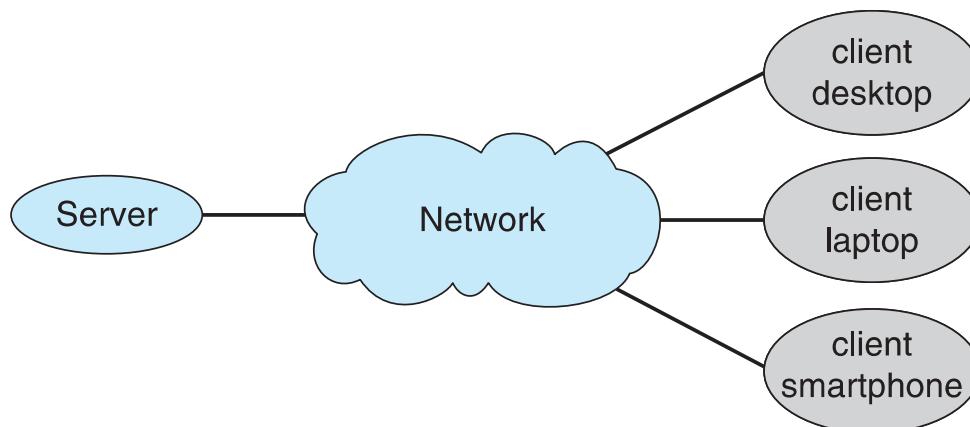
- Handheld **smartphones**, **tablets**, etc
- What is the **functional difference** between them and a **traditional laptop**?
 - **Extra feature** – more OS features: **GPS**, **gyroscope** (陀螺儀)
 - Allows new types of **apps** like ***augmented reality***
 - Use **IEEE 802.11 wireless**, or **cellular (3G or 4G) data networks** for connectivity
- Leaders are **Apple iOS** and **Google Android**





Computing Environments – Client-Server

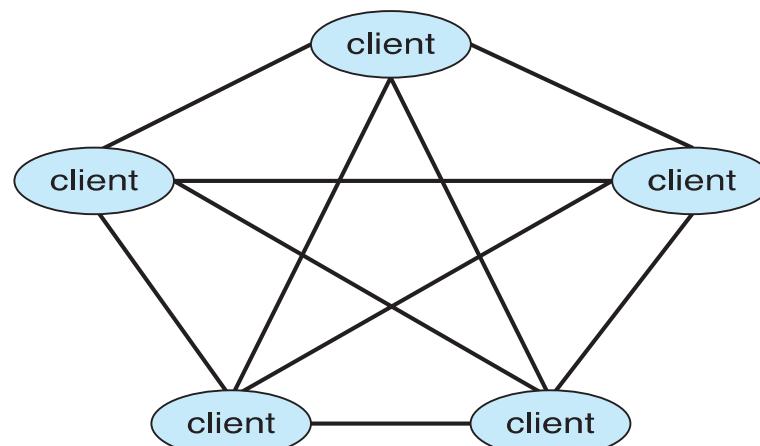
- Client-Server Computing
 - Dumb terminals replaced by smart PCs
- Many systems now **servers**, responding to requests generated by **clients**
 - **Compute-server system** provides an **interface** to **client** to request **services** (e.g. database)
 - **File-server system** provides **interface** for **clients** to store and retrieve **files**





Computing Environments - Peer-to-Peer

- Another model of **distributed system**
- **P2P** does not distinguish **clients** and **servers**
 - Instead **all nodes** are considered **peers**
 - May each act as **client**, **server** or both
 - Node must join **P2P network**
 - ▶ Registers its service with **central lookup service** on network, or
 - ▶ Broadcast request for service and respond to requests for service via **discovery protocol**
 - Examples include **Napster** and **Gnutella**, **Voice over IP (VoIP)** such as **Skype** or **Line**





Computing Environments - Virtualization

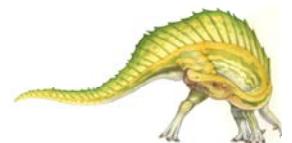
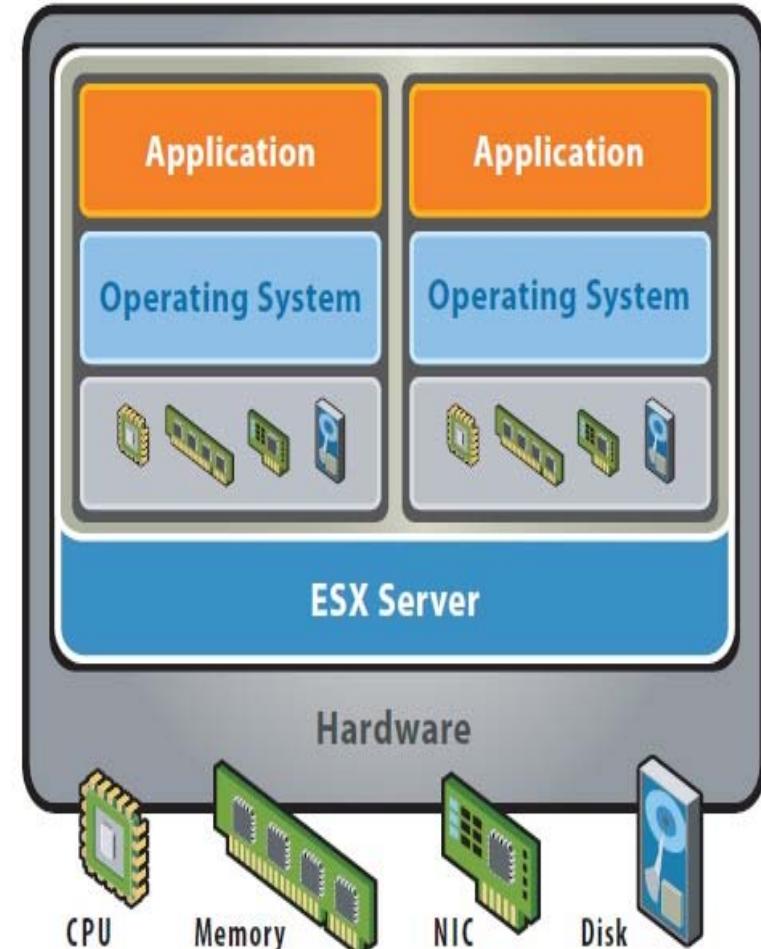
- Allows **operating systems** to run **applications** within **other OSes**
 - Vast and growing industry
- **Emulation** used when **source CPU type** different from **target type**
(e.g. **PowerPC** to Intel **x86**)
 - Generally **slowest** method
 - When **computer language** not compiled to **native code** –
Interpretation
- **Virtualization** – running **guest OSes** natively compiled for **CPU**
 - Consider **VMware** running **WinXP guests**, each running applications, all on native **Win-7 host OS**
 - **VMM** provides virtualization services





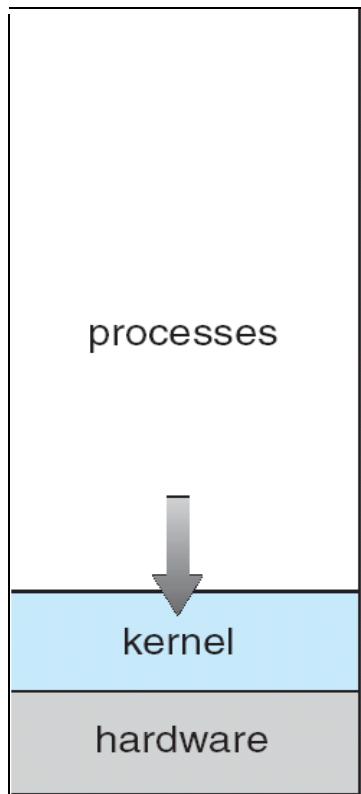
Computing Environments - Virtualization

- Laptops or desktops running multiple OSes for **exploration** or **compatibility**
 - Apple laptop running **Mac OS X** host, **Windows** as a guest
 - **Developing apps for multiple OSes** without having multiple systems
 - **QA testing applications** without having multiple systems
 - **Executing and managing compute environments** within **data centers**
- **VMM** (virtual machine monitor) can run natively, and they are the **host OS**
 - VMware ESX and Citrix XenServer



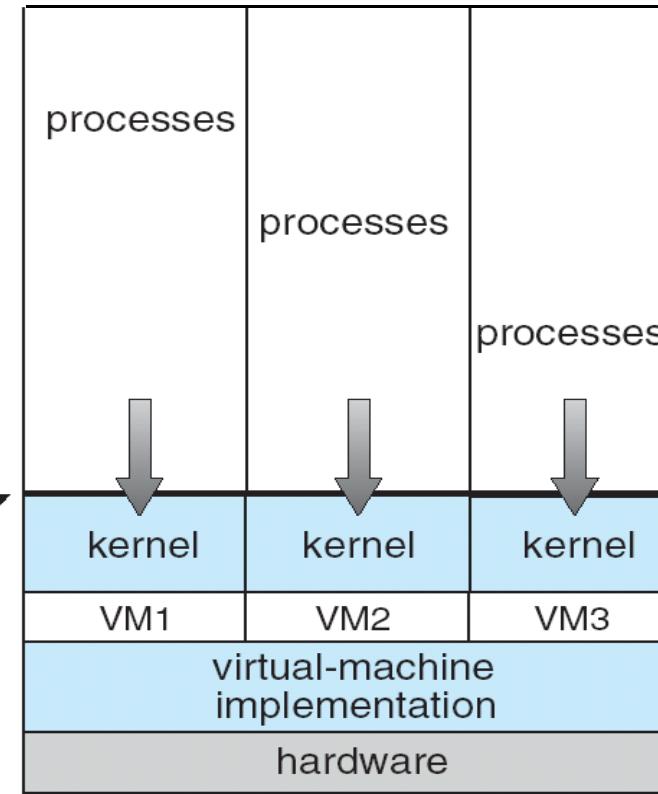


Computing Environments - Virtualization



(a)

(a) **Non-virtual** machine



(b)

(b) **virtual** machine





Computing Environments – Cloud Computing

- Delivers **computing**, **storage**, even **apps** as a **service** across a **network**
- **Logical extension of virtualization** as based on **virtualization**
 - **Amazon EC2** has **thousands of servers**, **millions of VMs**, PBs of storage available across the Internet, pay based on usage
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - **Software as a Service (SaaS)** – one or more **applications** available via the Internet (e.g. word processor)
 - **Platform as a Service (PaaS)** – **software stack** ready for application use via the Internet (e.g. a **database server**)
 - **Infrastructure as a Service (IaaS)** – **servers or storage** available over Internet (e.g. storage available for backup use)





Open-Source Operating Systems

- **Operating systems** made available in **source-code** format rather than just **binary closed-source** (早期, 1950, 開放原碼盛行)
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “**copyleft**” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
 - **Linus Torvalds** (a student in Finland), released a UNIX-like kernel in **1991**.
- Can use **VMM** like **VMware Player** (Free on Windows), **Virtualbox** (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run **guest operating systems** for exploration

