



Chapter 3: Process Concept

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems





Process Concept


- An **operating system** executes a variety of **programs**:
 - **Batch system** – **jobs**
 - **Time-shared systems** – **user programs** or **tasks**
- Textbook uses the terms **job** and **process** almost interchangeably
- **Program** is **passive** entity stored on disk (**executable file**), **process** is **active**
 - **Program** becomes **process** when executable file loaded into memory
 - **Process** – a program in execution
- Execution of **program** **started via** ***GUI mouse clicks, command line entry of its name***, etc
- A **program** (e.g. **notepad.exe**) may be **executed several times** =>
 - different **processes** run currently
 - each **process** has its own **data**, **stack**, and **heap**





Process in Memory


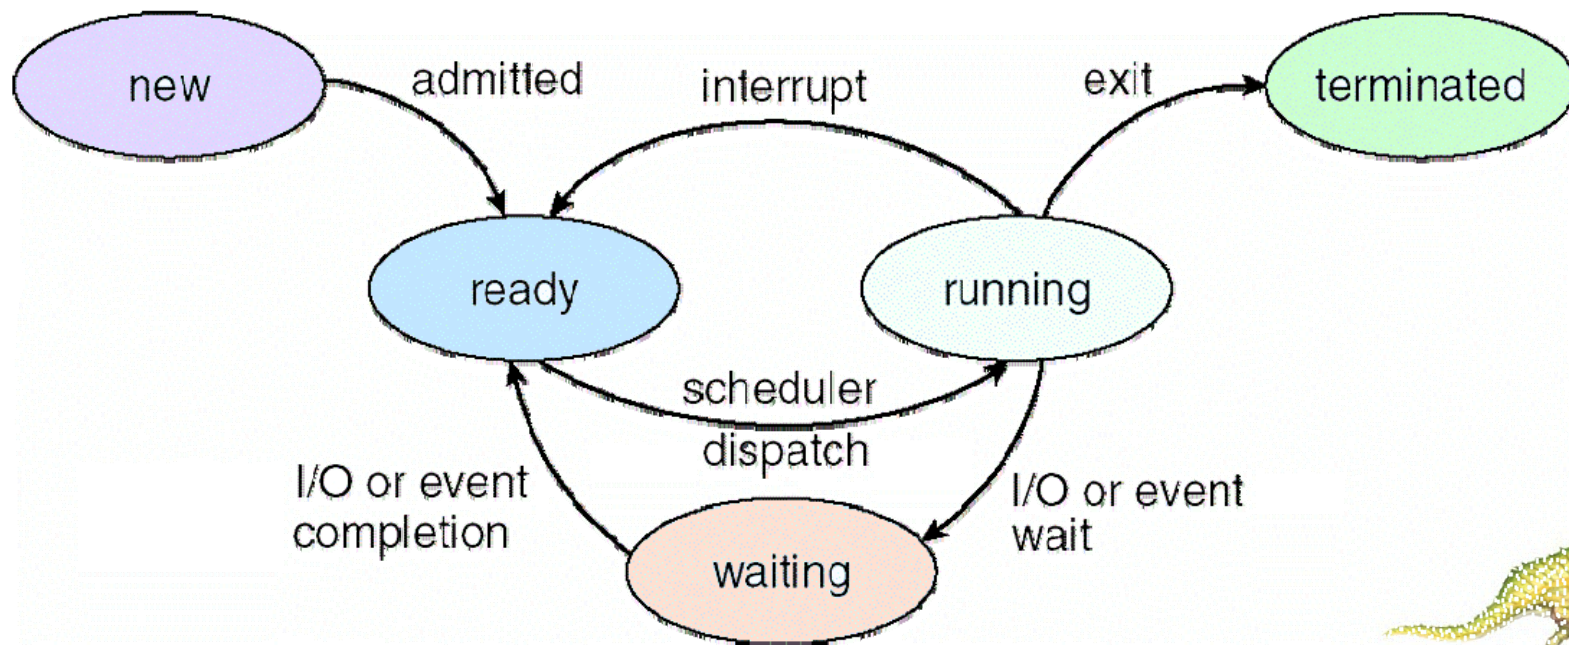
- A **process** includes:
 - **program counter**, **registers**
 - **stack**
 - ▶ containing *function parameters*, *return addresses*, *automatic variable*
 - **heap**
 - ▶ containing *memory dynamically allocated* during run time
 - **data** (*static variables*, *constants*)
 - **text** (*object code*)
- A **object program file** just includes **data** and **text** sections





Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: **Instructions** (object code) are being executed (occupy **processor**)
 - **waiting**: The process is waiting for **some event** to occur
 - **ready**: The process is **waiting to be assigned to a processor**
 - **terminated**: The process has **finished execution**

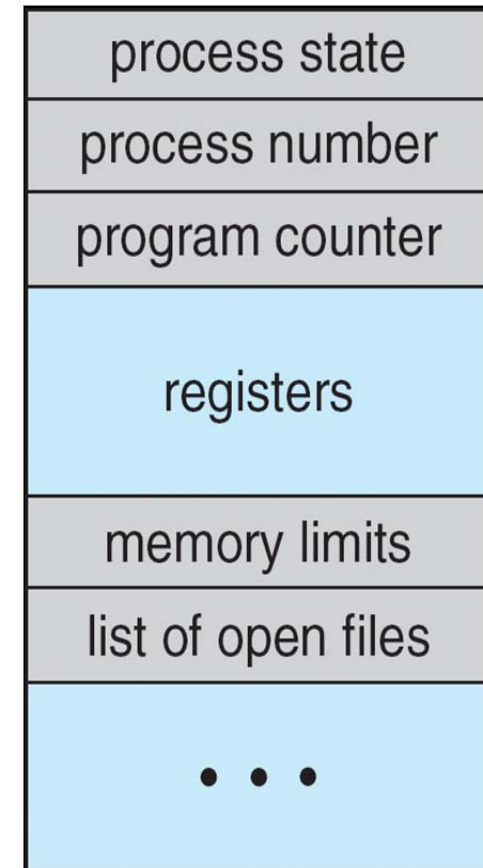




Process Control Block (PCB)

Information associated with each process
(also called **task control block**)

- **Process state** – running, waiting, etc
- **Program counter** – location of instruction to next execute
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling information** - priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files



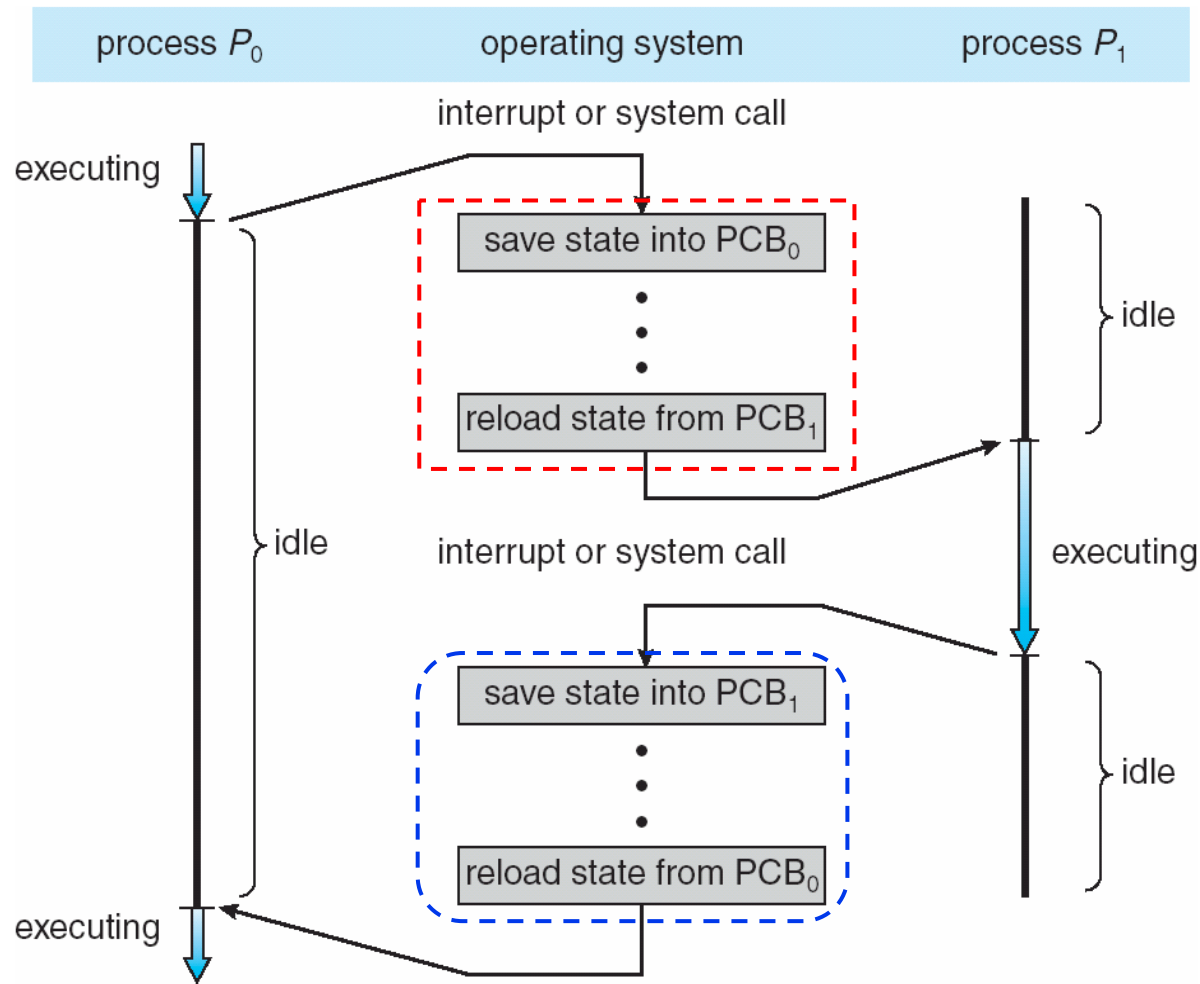
pid (process id)

base, limit registers





CPU Switch From Process to Process






Process Scheduling

- **Process scheduler** selects among available processes for next execution on CPU
- Maintains **scheduling queues** of processes. Processes migrate among the various queues
 - **Job queue** – set of all jobs in the system
 - ▶ **batch** : executes commands when **system load level** permit, e.g. the **load average** drops below 0.8
 - **Ready queue** – set of all processes residing in main memory, **ready** and waiting to execute (**wait CPU**)
 - **Device queues** – set of processes waiting for an **I/O device**
 - **Event queues** – waiting for an event to occur, e.g. **timer**, **semaphore**



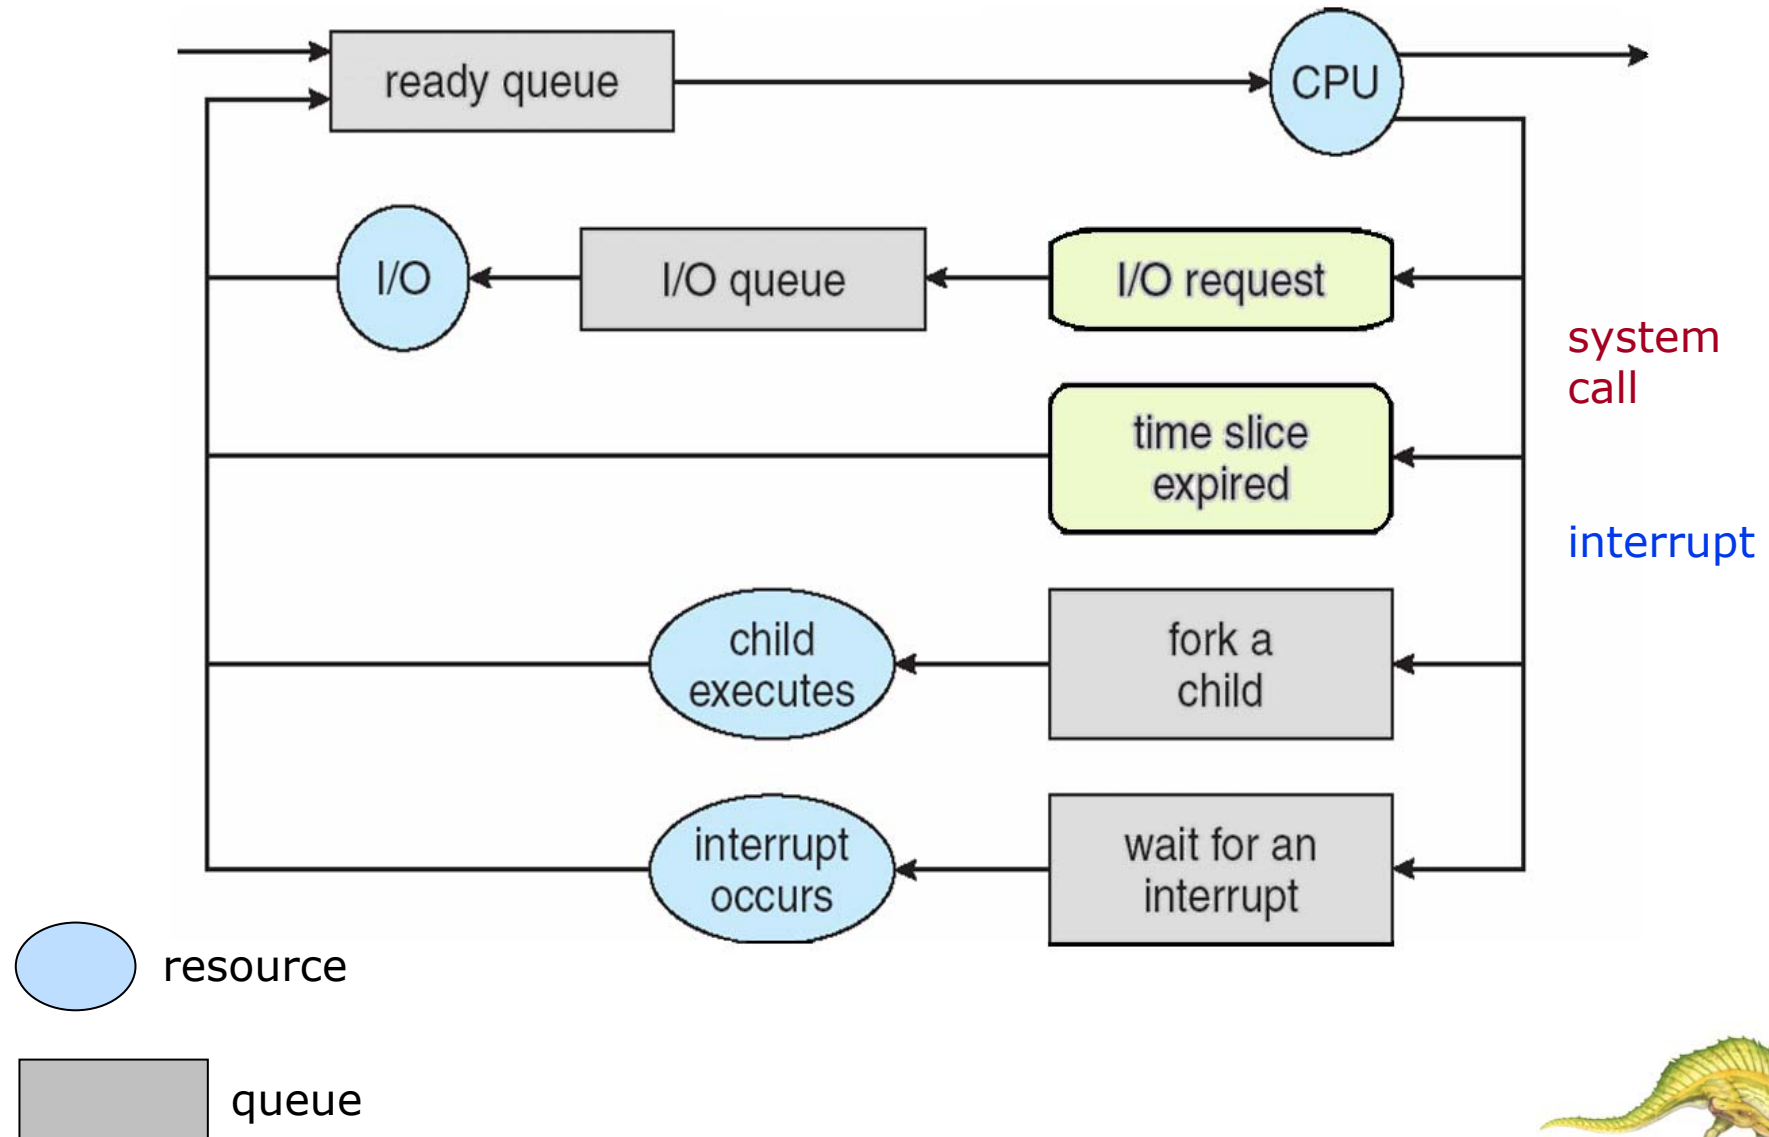


Ready Queue & I/O Device Queues





Queueing-diagram of Process





Schedulers

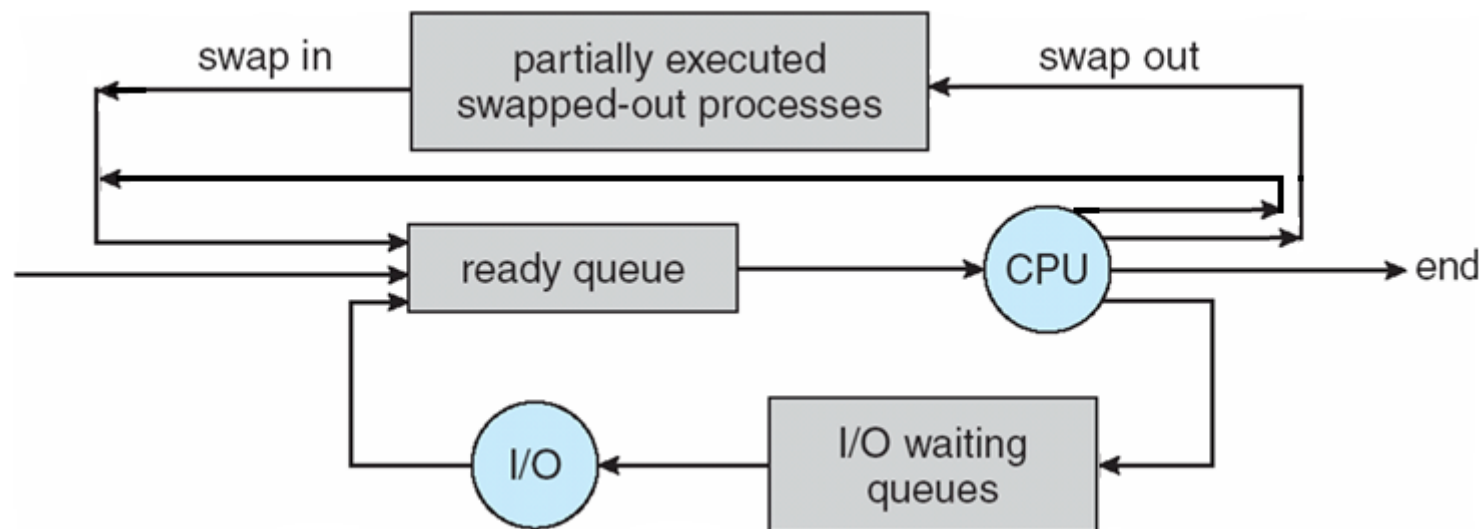
- **Long-term scheduler** (or **job** scheduler)
 - **selects** which **processes** should be brought into the **ready queue**
 - controls the **degree of multiprogramming** (but by user in MS windows)
- **Short-term scheduler** (or **CPU** scheduler)
 - **selects** which process should be executed next and **allocates CPU**
 - **invoked very frequently => must be fast.**
 - ▶ If it take 10 ms for 100 ms, then $10 / (10+100) = 9\%$





Medium Term Scheduling

- **Processes** can be described as either:
 - **I/O-bound process** – spends more time **doing I/O** than computations, many short CPU bursts
 - **CPU-bound process** – spends more time **doing computations**; few very long CPU bursts
- **Remove process** from memory, **store on disk**, **bring back in** from disk to continue execution (**swapping**)
- To **improve the process mix** in some **timesharing systems**





Context Switch

- When CPU switches to another process, the system must
 - **save** the state of the old process and
 - **load** the saved state for the new process via a **context switch**
- Context of a process represented in the **PCB**
- Context-switch time is overhead; the system does no useful work while switching
 - The more **complex** the **OS** and the **PCB** => longer the **context switch**
- Time dependent on hardware support
 - Some hardware provides **multiple sets of registers** per CPU => multiple **contexts** loaded at once





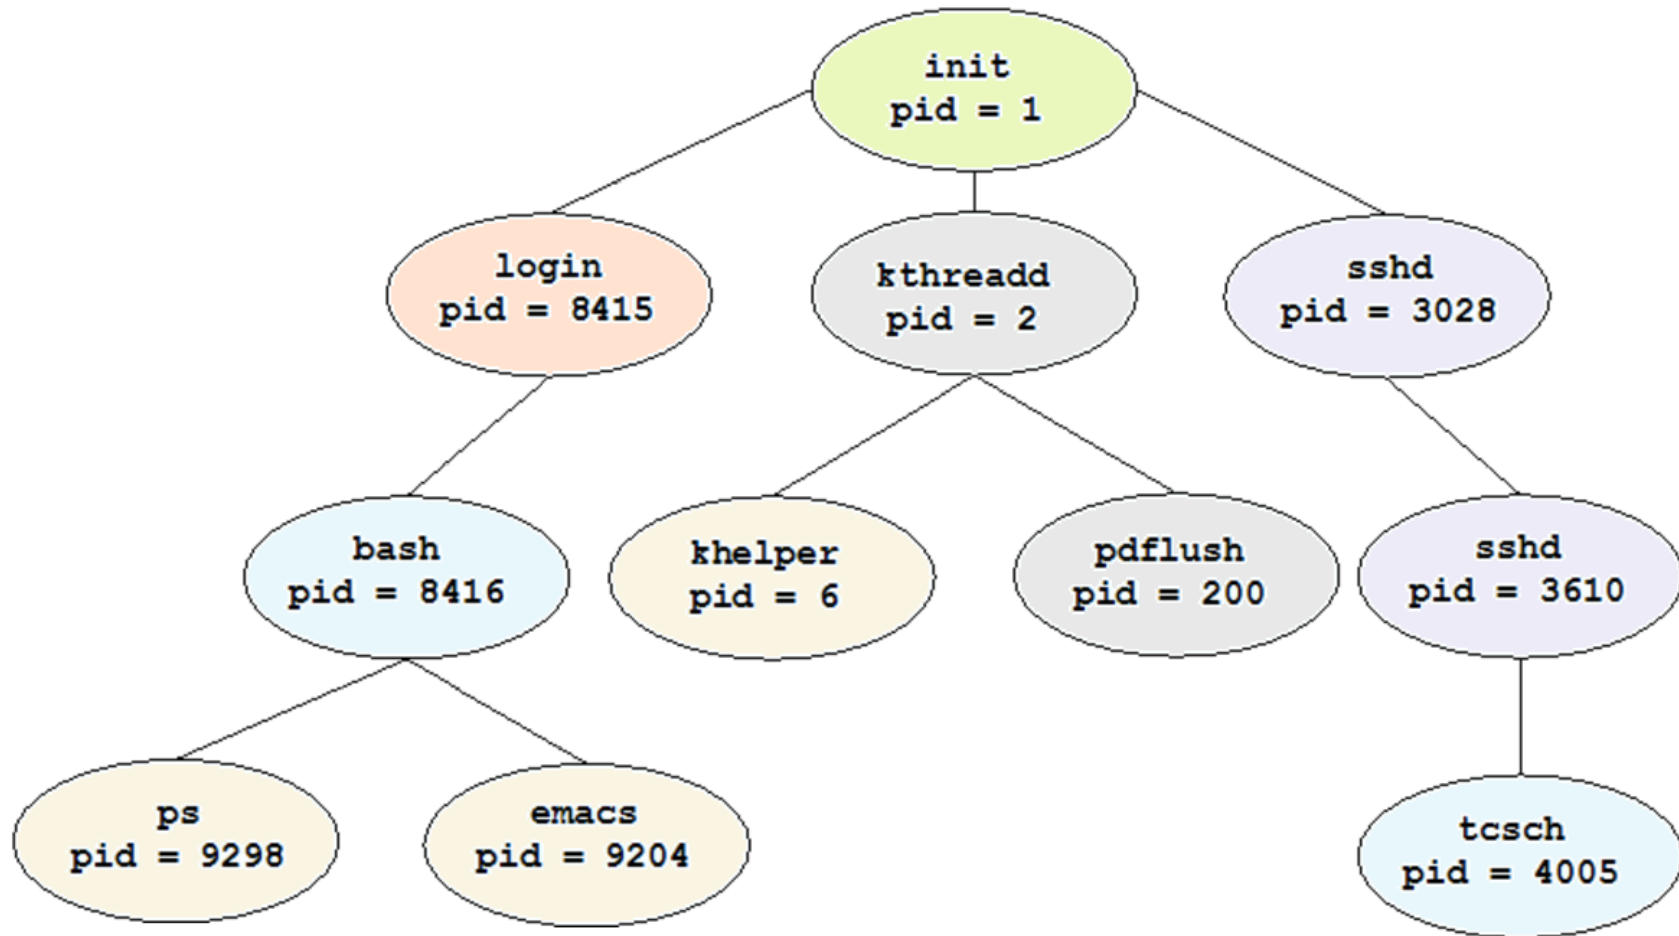
Process Creation (1)

- **Parent** process create **children** processes, which, in turn create other processes, forming a tree of processes
- **Process** identified and managed via a process identifier (**pid**)
- **Resource sharing**
 - Parent and children **share all resources**
 - Children **share subset** of parent's resources
 - Parent and child **share no resources**
- **Execution**
 - Parent and children **execute concurrently**
 - Parent **waits** until children **terminate**





A tree of processes on a Linux





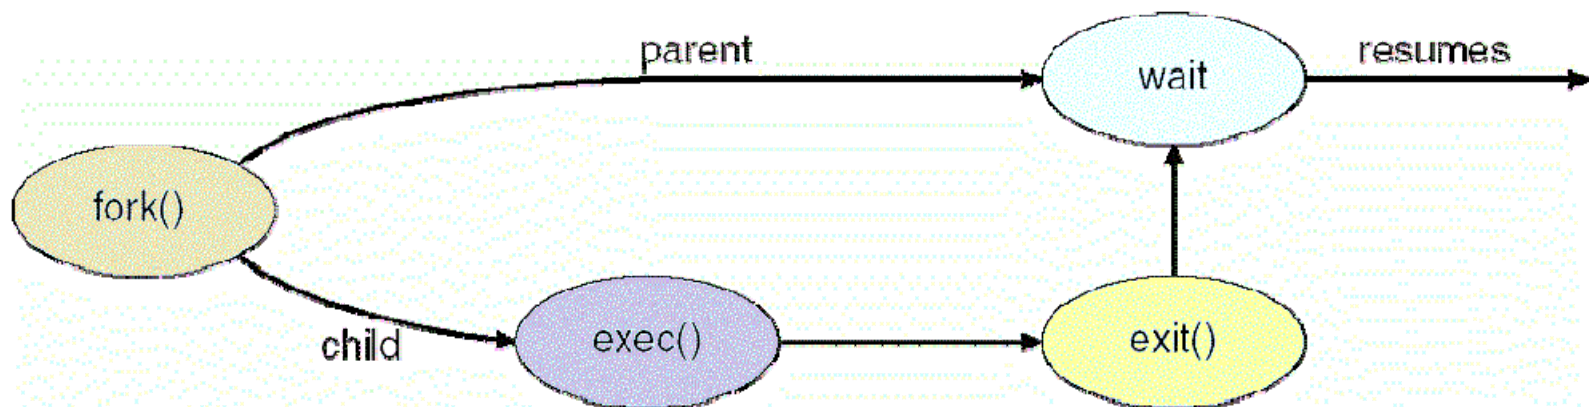
Process Creation (2)

■ Address space

- Child **duplicate** of parent
- Child has a **program loaded** into it

■ UNIX examples

- **fork** system call **creates** new process
- **exec** system call used after a **fork** to **replace** the process' memory space with a **new program**





C Program Forking Separate Process

```
#include <unistd.h>
int main()
{
    pid_t pid;

    /* fork another process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Parent process

Child process





Creating Separate Process Using Win32 API

```
#include <windows.h>
#include <stdio.h>

int main( VOID )
{
    STARTUPINFO si;

    //-----
    // create child process
    if( !CreateProcess( NULL,    //use command line
        "C:\\\\WINDOWS\\system32\\mspaint.exe", // Command line
        NULL,                // don't inherit process handle
        NULL,                // don't inherit thread handle
        FALSE,               // disable handle inheritance
        0,                   // No creation flags
        NULL,                // Use parent's environment block
        NULL,                // Use parent's starting directory
        &si,
        &pi )
    )
    //-----
    {
        fprintf( stderr,
            "Create Process Failed." );
        return -1;
    }

    // parent will wait for
    // the child to complete

    WaitForSingleObject(
        pi.hProcess, INFINITE );
    printf( "Child Complete" );

    // close handles
    CloseHandle( pi.hProcess );
    CloseHandle( pi.hThread );
}
```





Process Termination

- **Process** executes last statement and asks the **OS** to **delete** it, e.g. `exit(1)`
 - Output data from child to parent (via `wait()`)
 - ▶ `pid = wait(&status);`
 - Process' **resources** are **deallocated** by operating system
- **Parent** may **terminate children processes**, e.g. `kill(pid, signal)`
 - Child has **exceeded allocated resources**
 - Task assigned to child is **no longer required**
 - **If parent is exiting**
 - ▶ Some OS do not allow child to continue if its parent terminates
 - All children terminated -- **cascading termination**
- **If no parent waiting**, then terminated process is a **zombie**
 - Its entry in **process table** must remain there until the parent call `wait()`
- **If parent terminated**, **process** is **orphan** (assigning *init* as new parent)
 - e.g., `nohup abc &`





Multiprocess Architecture – Chrome Browser

- Many **web browsers** ran as **single process**
 - If one web site causes trouble, entire browser can **hang** or **crash**
- **Google Chrome Browser** is **multiprocess** with 3 categories
 - **Browser** process manages **user interface, disk and network I/O**
 - **Renderer** process renders **web pages, deals with HTML, Javascript, new one for each website opened**
 - ▶ Run in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in





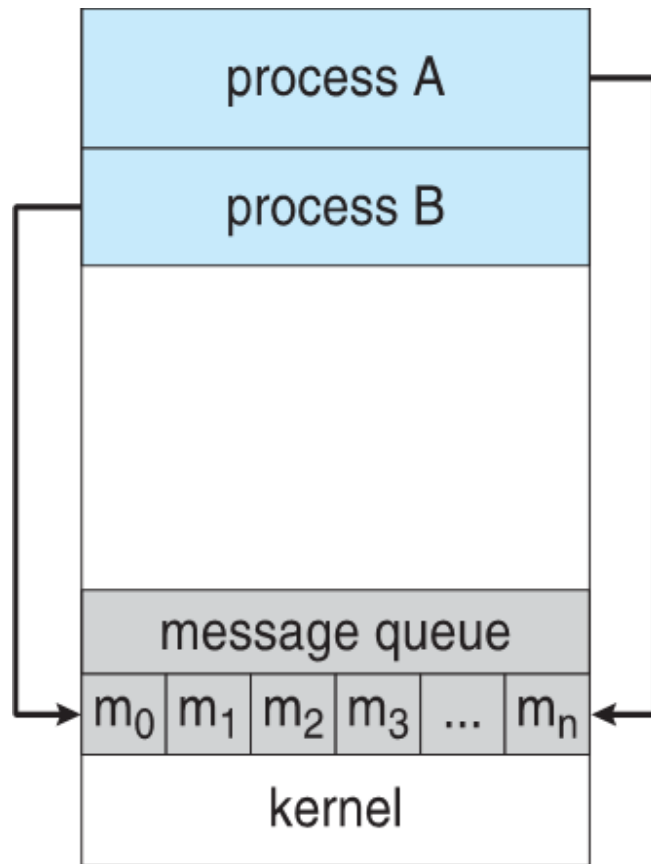
Interprocess Communication

- **Processes** within a system may be **independent** or **cooperating**
- **Reasons** for **cooperating processes**:
 - Information sharing
 - Computation speedup
 - **Modularity** (e.g. `ls -l | grep abc`)
 - Convenience
- **Cooperating process** can **affect** or be **affected** by other processes, including **sharing data**
- **Cooperating processes** need **interprocess communication (IPC)**
- **Models of IPC**
 - **Shared memory**
 - **Message passing**






Communications Models



(a)

message passing



(b)

shared memory





Producer-Consumer Problem

■ Paradigm for cooperating processes

- **producer** process produces information that is consumed by a **consumer** process
- **bounded-buffer** assumes that there is a **fixed buffer size**

■ One Solution : **Bounded-Buffer** on **Shared-Memory**

```
/* Declaration of the shared data */  
#define BUFFER_SIZE 10  
typedef struct {  
    ...  
} item;  
  
struct shrd {  
    item buffer[BUFFER_SIZE];  
    int in = 0;  
    int out = 0;  
} *mp;
```

`mp = shmat(sizeof(shrd));`
`mp->in ++;`





Bounded-Buffer Producer

- Solution is correct, but can only use BUFFER_SIZE - 1 elements

```
item nextProduced ;
while (true) {
    /* Produce an item in nextProduced */

    while ( ( (in + 1) % BUFFER_SIZE) == out )
        ; /* do nothing -- no free buffer */

    buffer[ in ] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Bounded-Buffer Consumer

```
item nextConsumed ;
while (true) {
    while (in == out)
        ; // do nothing

    nextConsumed = buffer[ out ];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in nextConsumed */
}
```





IPC – Message Passing

- **Mechanism** for processes to **communicate** and to **synchronize** their actions
- **Message system** – processes communicate with each other **without** resorting to **shared variables**
- **Message passing** facility provides two operations:
 - **send**(*message*) , message size **fixed** or **variable**
 - **receive**(*message*)
- If *P* and *Q* wish to communicate, they need to establish a **communication link**, and exchange messages via **send/receive**
- **Implementation** of **communication link**
 - **direct** or **indirect** communication, i.e. **naming**
 - **Synchronous** or **asynchronous** communication
 - automatic or explicit **buffering**





Direct Communication

- Processes must **name** each other explicitly:
 - **send** (P , *message*) – send a message to process P
 - **receive**(Q , *message*) – receive a message from process Q
 - If the name of P or Q changed, all such names must be changed
- Properties of **communication link**
 - Links are established **automatically**
 - A link is associated with **exactly one pair** of communicating processes
 - The link may be **unidirectional**, but is usually **bi-directional**





Indirect Communication

- **Messages** are sent to and received from **mailboxes** (ports)
 - Each **mailbox** has a unique **id**
 - **Processes** can communicate only if they share a **mailbox**
- **Primitives** are defined as:
 - send**(*A, message*) – send a message to **mailbox A**
 - receive**(*A, message*) – receive a message from **mailbox A**
- Properties of **communication link**
 - **Link** established only if processes share a common **mailbox**
 - A **link** may be associated with **many processes**
 - Each pair of processes may **share** several **communication links**





Indirect Communication

■ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 sends; P_2 and P_3 receive
- Who gets the message?

■ Solutions

- Allow a link to be associated with at most two processes
- Allow **only one process** at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Message Passing -- Synchronization

- Message passing may be either **blocking** or **non-blocking**
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
e.g. `SendMessage(hWnd, MsgID, WParam, LParam)`
 - **Blocking receive** has the receiver block until a message is available
e.g. `GetMessage(lpMsg, hWnd, FirstMsgID, LastMsgID)`
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
e.g. `PostMessage()`
 - **Non-blocking receive** has the receiver receive a valid message or null
e.g. `PeekMessage()`





Message Passing -- Buffering

- **Queue** of messages attached to the link; implemented in one of three ways
 1. **Zero capacity** – 0 messages
Sender must **wait** for **receiver** (**rendezvous**)
 2. **Bounded capacity** – **finite length** of **n** messages
Sender must **wait** if **link full**
 3. **Unbounded capacity** – **infinite length**
Sender **never waits**





Example: POSIX Shared Memory

- Process first **creates** shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment to share it

- **Set the size** of the object

```
ftruncate(shm_fd, 4096);
```

- **Map** the shared memory object

```
ptr = mmap(0, 4096, PROT_WRITE, MAP_SHARED, shm_fd, 0);
```

- Now the process could **write to** the shared memory

```
sprintf(ptr, "Writing to shared memory");
```





IPC POSIX Producer

```
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main()
{
    const int SIZE = 4096;
    const char *name = "OS";
    const char *message0= "Studying ", *message1= "Operating Systems ";
    const char *message2= "Is Fun!";
    int shm_fd;
    void *ptr;

    /* create the shared memory segment */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory segment */
    ftruncate(shm_fd, SIZE);

    /* now map the shared memory segment in the address space of the process */
    ptr = mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);
    if (ptr == MAP_FAILED) {
        printf("Map failed\n");    return -1;
    }

    /* Now write to the shared memory region.
     * Note we must increment the value of ptr after each write. */
    sprintf(ptr, "%s", message0);    ptr += strlen(message0);
    sprintf(ptr, "%s", message1);    ptr += strlen(message1);
    sprintf(ptr, "%s", message2);    ptr += strlen(message2);
    return 0;
}
```





IPC POSIX Consumer

```
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main()
{
    const char *name = "OS";
    const int SIZE = 4096;
    int shm_fd;
    void *ptr;
    int i;

    /* open the shared memory segment */
    shm_fd = shm_open(name, O_RDONLY, 0666);
    if (shm_fd == -1) {
        printf("shared memory failed\n");
        exit(-1);
    }

    /* now map shared memory segment in the address space of the process */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);
    if (ptr == MAP_FAILED) {
        printf("Map failed\n");
        exit(-1);
    }


    /* now read from the shared memory region */
    printf("%s", ptr);
    /* remove the shared memory segment */
    if (shm_unlink(name) == -1) {
        printf("Error removing %s\n", name);
        exit(-1);
    }
    return 0;
}
```





Example: Windows Message Passing

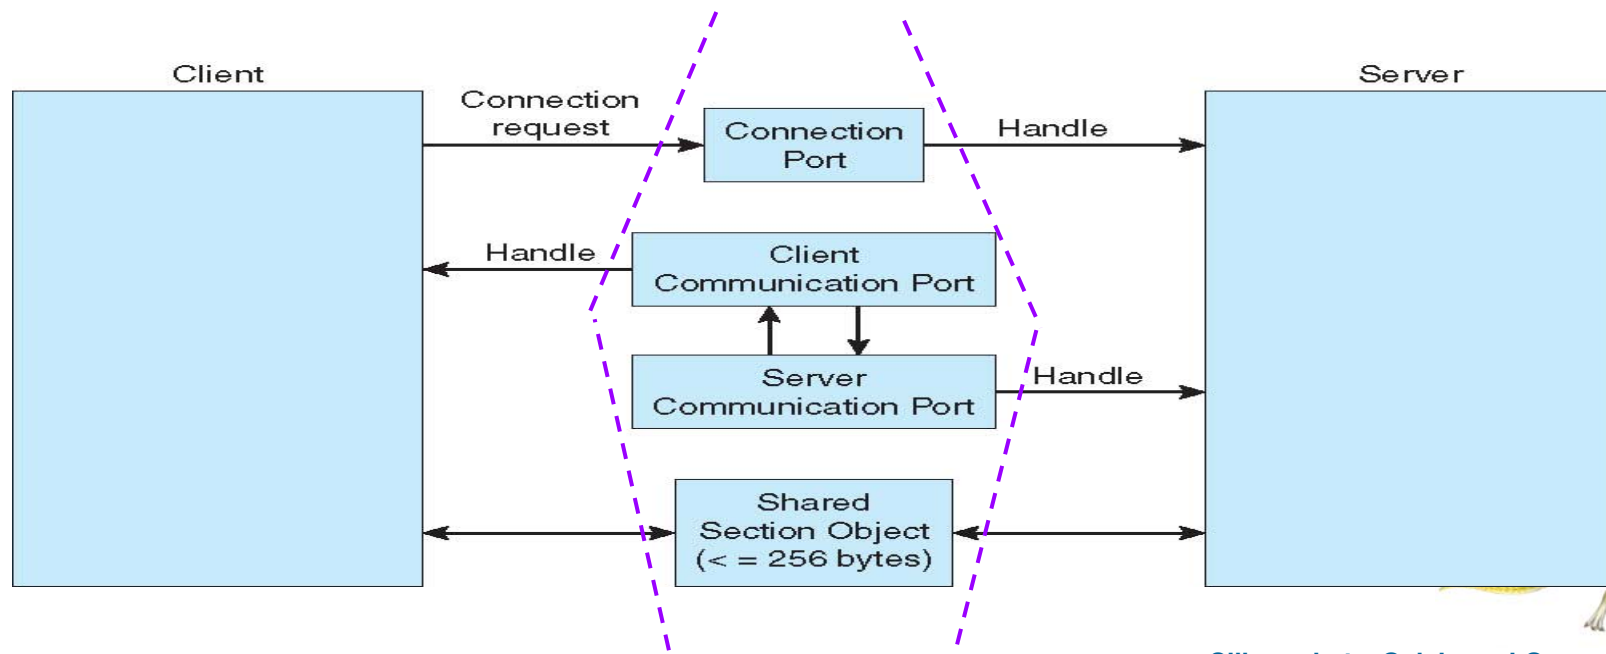
- **Windows** provides support for **multiple subsystems** (e.g. **Win32** server, **POSIX** server) with which application programs communicate via **message passing**
- **Message-passing** facility is called **advanced local procedure call (ALPC)**
 - Only works **between processes on the same system**
 - Uses **ports** (like **mailboxes**) to establish and maintain communication channels
- The **ALPC** facility in Windows is not part of the **Win32 API** (i.e. not visible to application programs)





Local Procedure Calls in Windows

- **Communication** works as follows:
 - The **client opens** a **handle** to the subsystem's **connection port object**
 - The **client** sends a **connection request**
 - The **server creates** two **private communication ports** and returns the handle to one of them to the client
 - The **client** and **server** use the corresponding **port handle** to **send messages** or **callbacks** and to **listen** for replies





Communications in Client-Server Systems


- Sockets
- Remote Procedure Calls
- Pipes





Socket Communication

- A **socket** is defined as an *endpoint for communication*
- Concatenation of **IP** address and **port**
- The **socket 161.25.19.8:1625** refers to **port 1625** on host **161.25.19.8**
- Communication consists between a pair of sockets





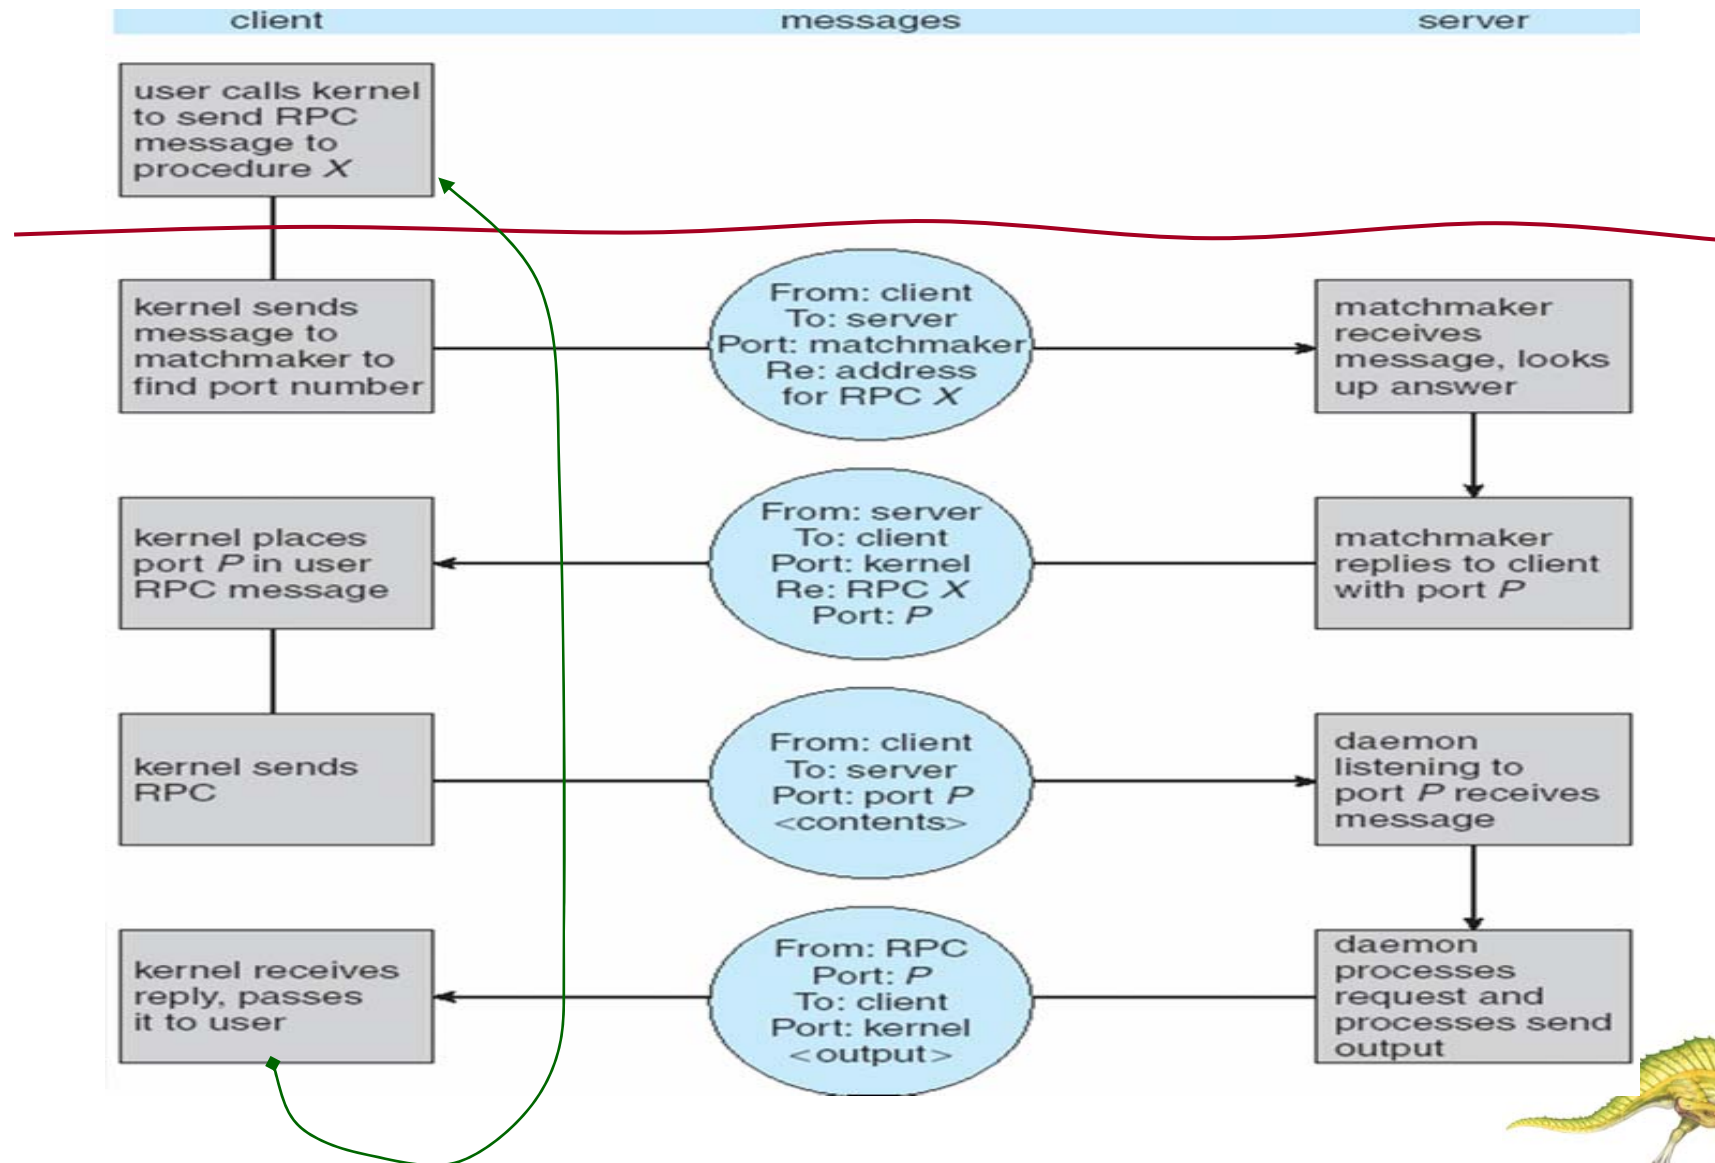
Remote Procedure Calls

- **Remote procedure call (RPC)**
 - abstract procedure calls between **processes** on **networked systems**
 - **Stubs** – **client-side proxy** for the actual **procedure** on the server
- The **client-side stub** **locates** the **server** and **marshalls** the **parameters**
- The **server-side stub** **receives** this message, **unpacks** the marshalled **parameters**, and **performs** the **procedure** on the server






Execution of RPC





Pipes

- A **pipe** acts as a conduit allowing two processes to communicate. One of the first **IPC mechanisms** in early **UNIX** systems.
- **Ordinary pipes** (e.g. `ls -l | grep abc`)
 - **unidirectional** -- **two pipes** must be used for **two-way communication**.
 - constructed using the **system call**, `pipe(int fd[])`
 - `fd[0]` is the descriptor for the **read-end**, `fd[1]` is for the **write-end**
 - Once the **processes terminate** the ordinary **pipe ceases to exist**
- **Named pipes (FIFOs)**
 - Once created, they **appear as typical files** in the **file system**, `mkfifo()`
 - **Continue to exist** after communicating processes have finished
 - Once established, **several processes can use it** for communication
- **UNIX** treats a **pipe** as **a special type of file**; thus, **pipes** can be accessed using ordinary system calls, `read()` and `write()`





Example: POSIX Pipe

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <string.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    pid = fork();    /* fork a child process */

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork failed");
        return 1;
    }
}
```

```
if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* write to the pipe */
    write(fd[WRITE_END], write_msg,
          strlen(write_msg)+1);

    /* close the write end of the pipe */
    close(fd[WRITE_END]);
}
else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);

    /* read from the pipe */
    read(fd[READ_END], read_msg,
          BUFFER_SIZE);
    printf("read %s", read_msg);

    /* close the write end of the pipe */
    close(fd[READ_END]);
}

return 0;
}
```

